

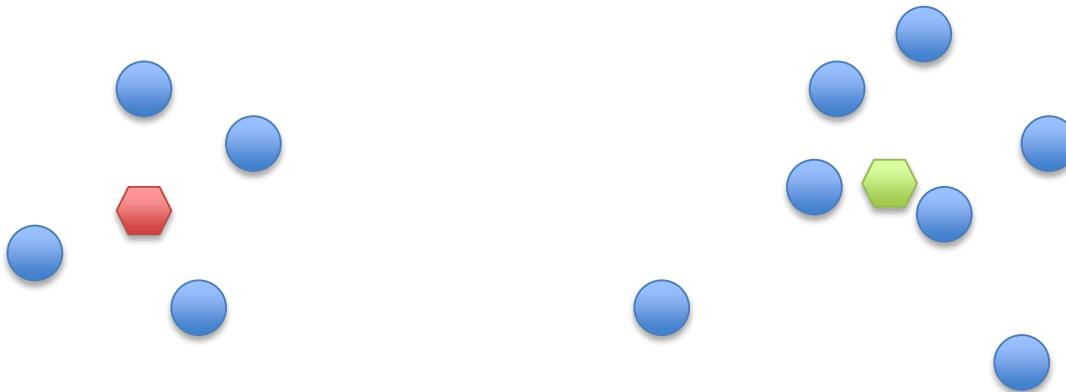
Spectral Clustering

Based on slides

(with thanks to Shannon Quinn,
William Cohen of Carnegie Mellon University, and J. Leskovec, A.
Rajaraman, and J. Ullman of Stanford University)

Partitional Clustering

- How do we partition a space to make the best clusters?
- Proximity to a cluster centroid.



Difficult Clusterings

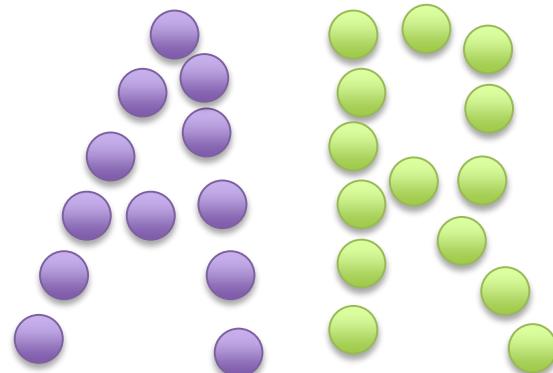
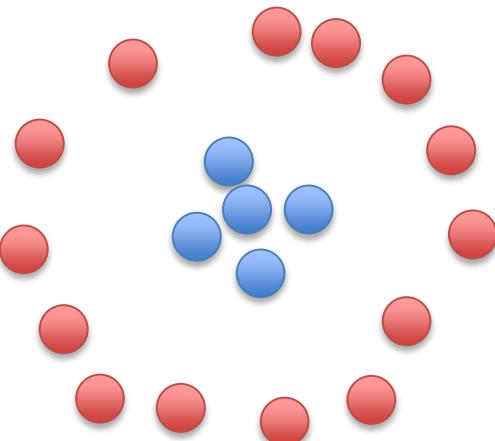
- But some clusterings don't lend themselves to a “centroid” based definition of a cluster.



- Spectral clustering allows us to address these sorts of clusters.

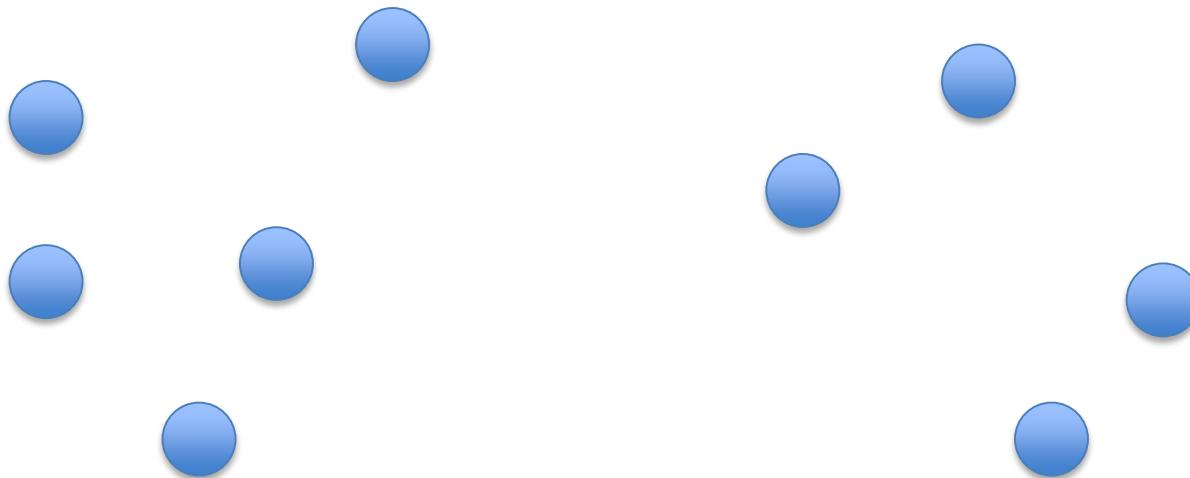
Difficult Clusterings

- These kinds of clusters are defined by points that are close **any** member in the cluster, rather than the **average** member of the cluster.



Graph Representation

- We can represent the relationships between data points in a graph.

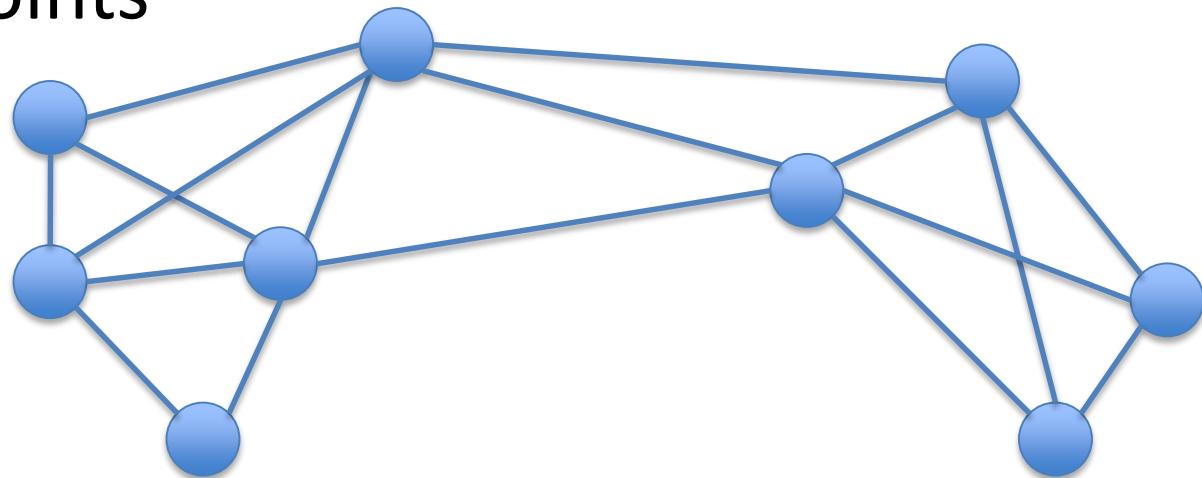


Graph definitions

- ε -neighborhood graph
 - Identify a threshold value, ε , and include edges if the affinity between two points is greater than ε .
- k-nearest neighbors
 - Insert edges between a node and its k-nearest neighbors.
 - Each node will be connected to (at least) k nodes.
- Fully connected
 - Insert an edge between every pair of nodes.

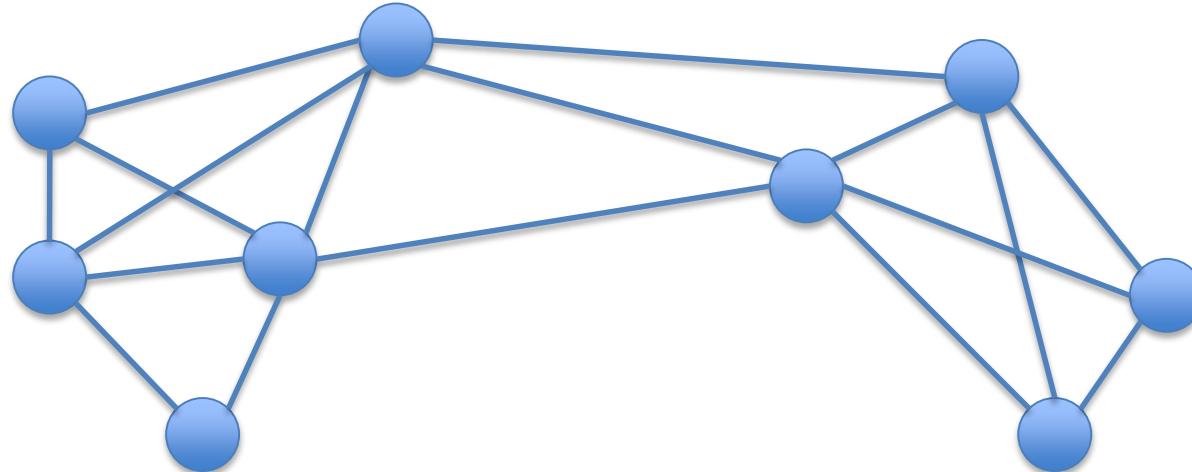
Graph Representation

- We can represent the relationships between data points in a graph.
- Weight the edges by the similarity between points



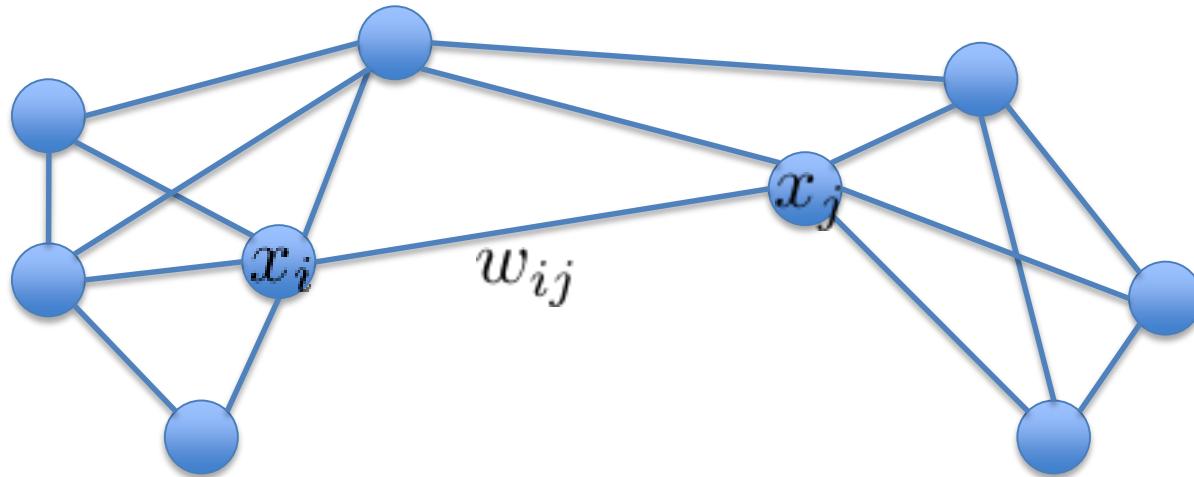
Representing data in a graph

- What is the best way to calculate similarity between two data points?
- Distance based: $d(x_i, x_j) = \exp \left\{ \frac{\|x_i - x_j\|}{\sigma^2} \right\}$



Graphs

- Nodes and Edges
- Edges can be directed or undirected
- Edges can have weights associated with them



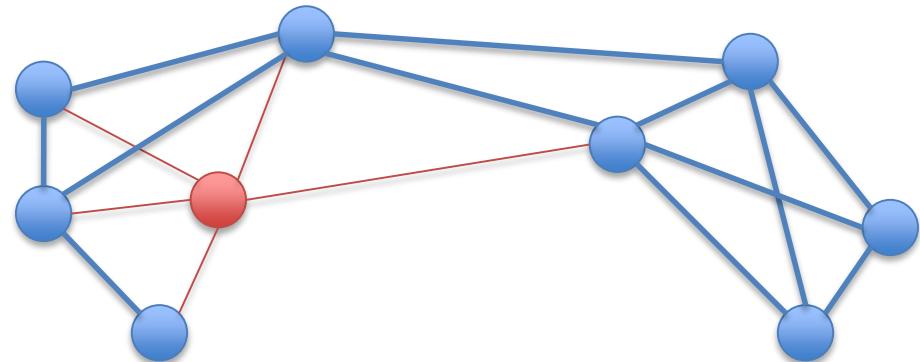
- Here the weights correspond to **pairwise affinity**

$$w_{ij} = d(x_i, x_j) = \exp \left\{ \frac{\|x_i - x_j\|}{\sigma^2} \right\}$$

Graphs

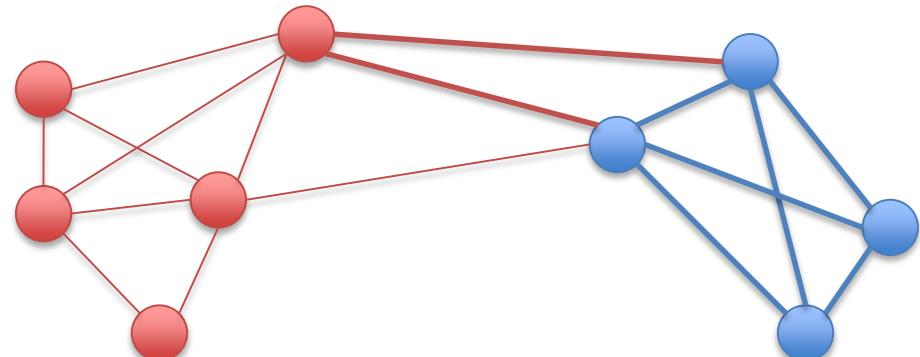
- Degree

$$D(x_i) = \sum_{j \in V} w_{ij}$$



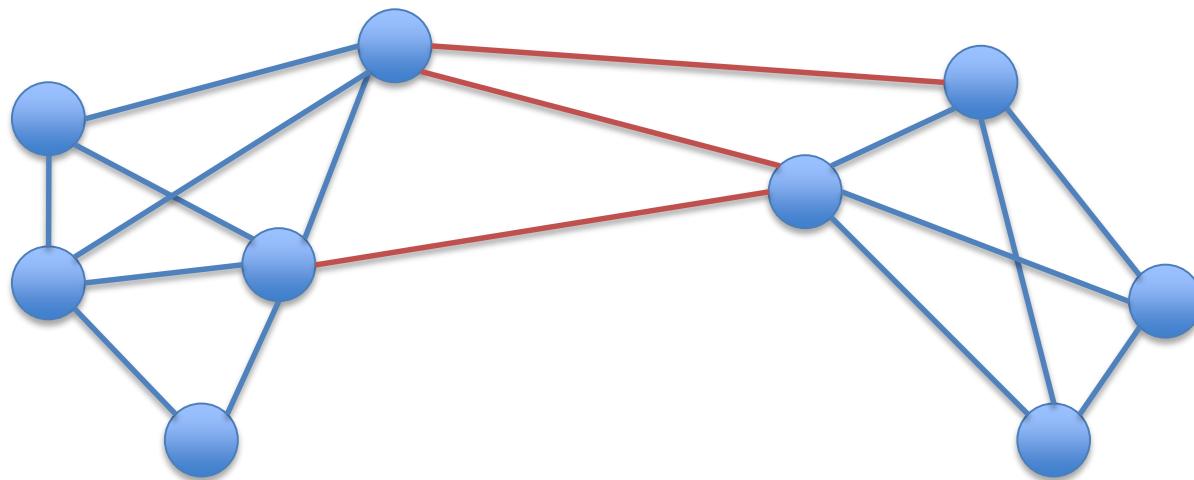
- Volume of a set

$$Vol(C) = \sum_{i \in C} D(x_i)$$



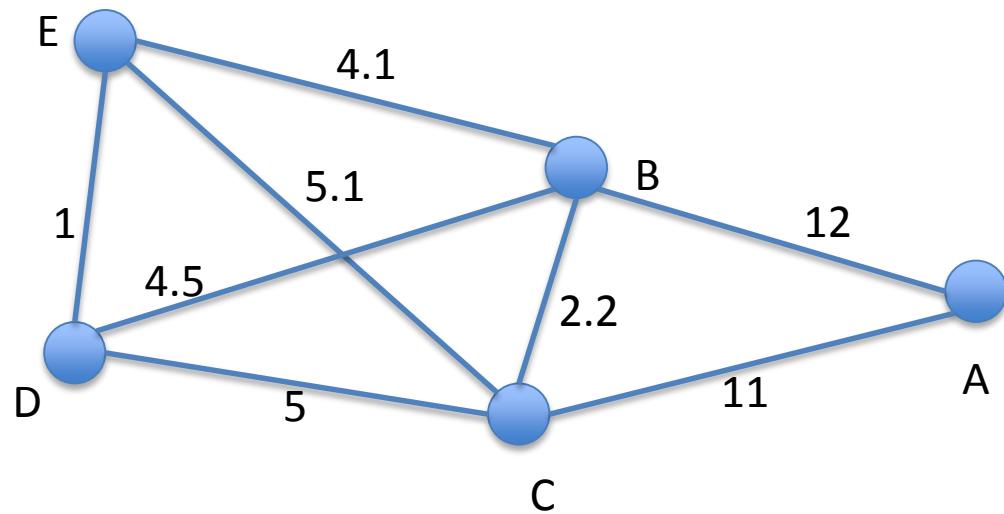
Graph Cuts

- The **cut** between two subgraphs is calculated as follows



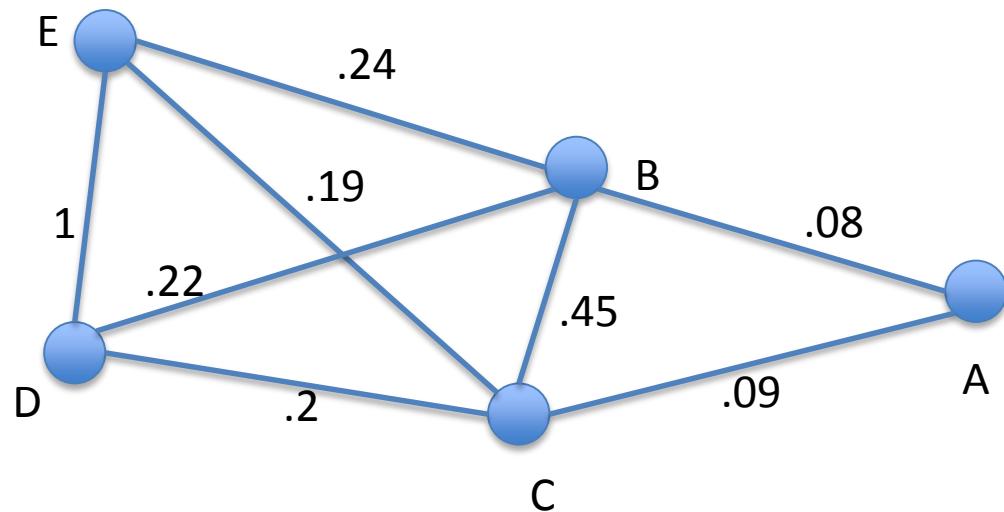
$$Cut(C_1, C_2) = \sum_{i \in C_1} \sum_{j \in C_2} w_{ij}$$

Graph Examples - Distance



Height	Weight
20	5
8	6
9	4
4	4
4	5

Graph Examples - Similarity



Height	Weight
20	5
8	6
9	4
4	4
4	5

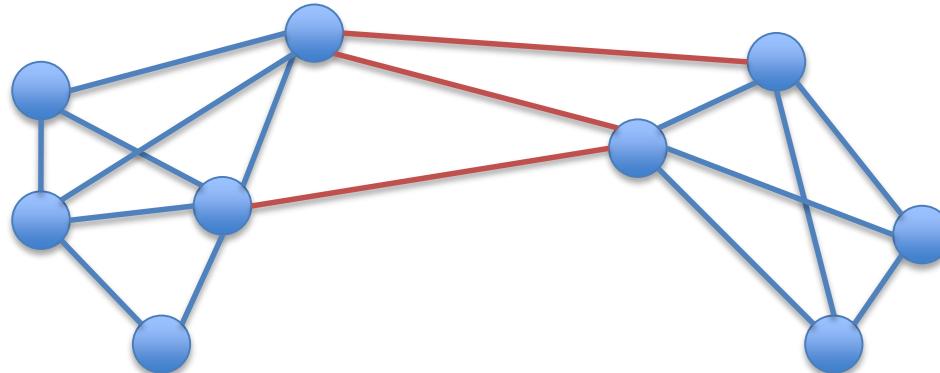
Intuition

- The minimum cut of a graph identifies an optimal partitioning of the data.
- Spectral Clustering
 - Recursively partition the data set
 - Identify the minimum cut
 - Remove edges
 - Repeat until k clusters are identified

Graph Cuts

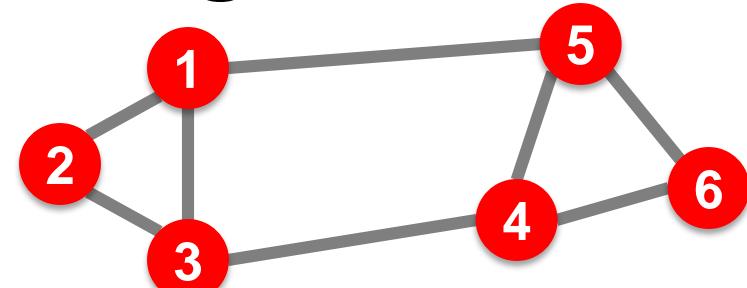
- Minimum (bipartitional) cut

$$\min Cut(C_1, C_2) = \sum_{i \in C_1} \sum_{j \in C_2} w_{ij}$$

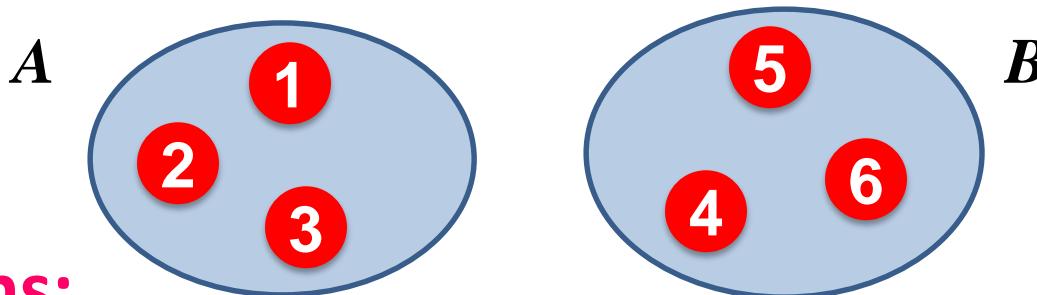


Graph Partitioning

- **Undirected graph**
- **Bi-partitioning task:**

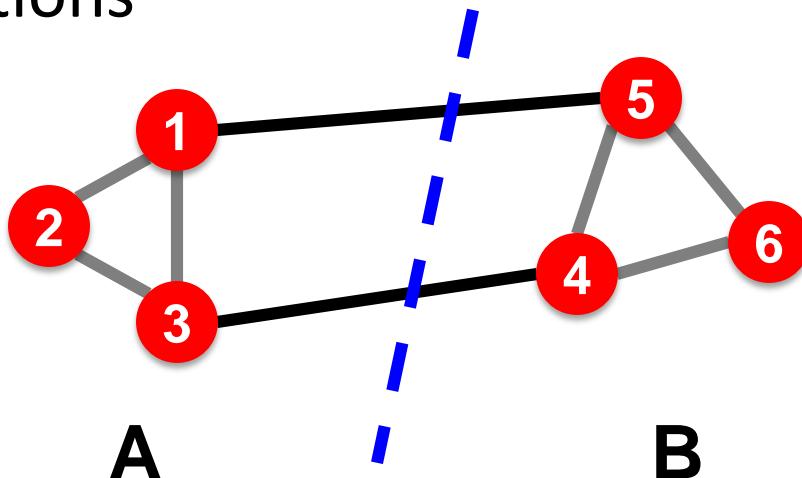


- Divide vertices into two disjoint groups
- **Questions:**
 - How can we define a “good” partition of ?
 - How can we efficiently identify such a partition?



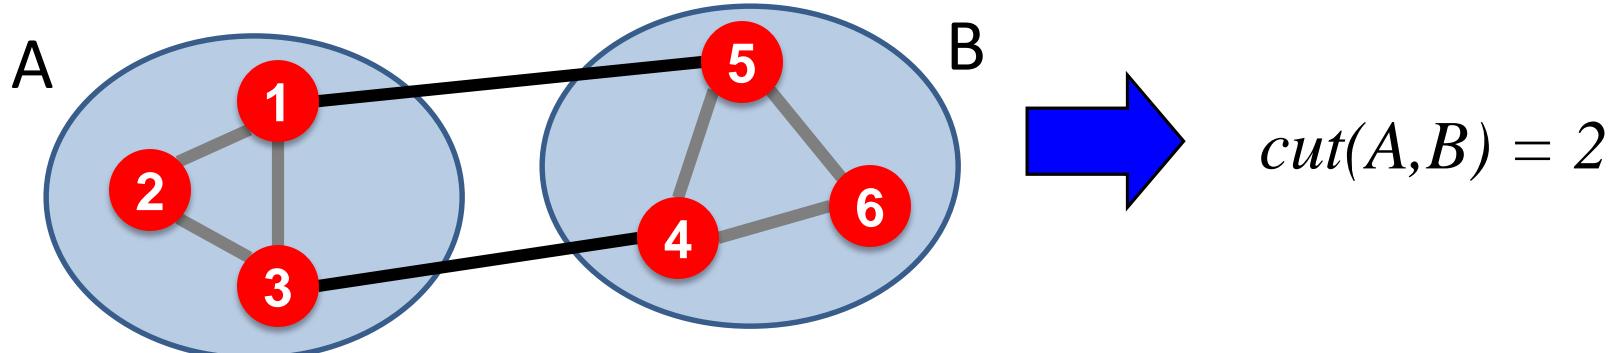
Graph Partitioning

- **What makes a good partition?**
 - Maximize the number of within-group connections
 - Minimize the number of between-group connections



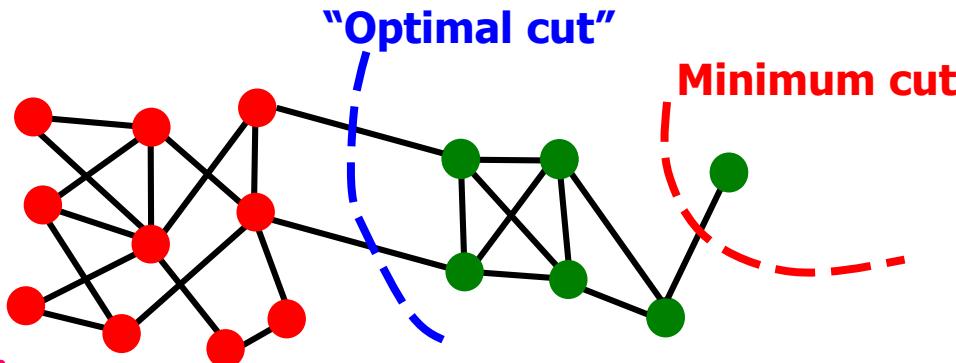
Graph Cuts

- Express partitioning objectives as a function of the “edge cut” of the partition
- Cut: Set of edges with only one vertex in a group:
$$cut(A, B) = \sum_{i \in A, j \in B} w_{ij}$$



Graph Cut Criterion

- **Criterion: Minimum-cut**
 - Minimize weight of connections between groups
$$\arg \min_{A,B} cut(A,B)$$
- **Degenerate case:**

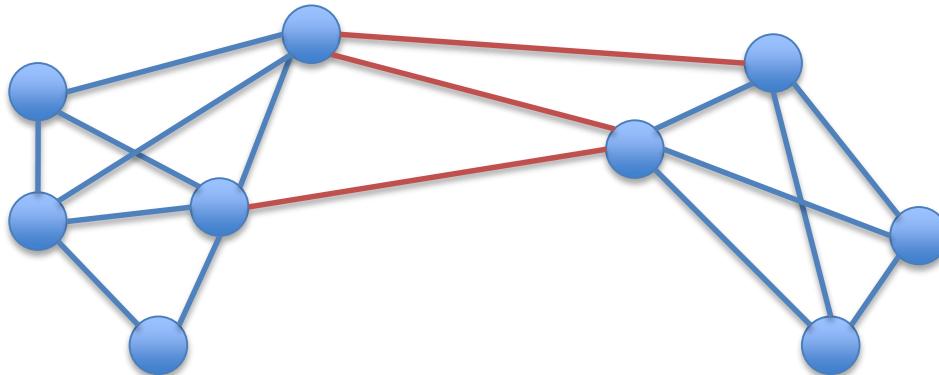


- **Problem:**
 - Only considers external cluster connections
 - Does not consider internal cluster connectivity

Graph Cuts

- Minimal (bipartitional) normalized cut.

$$\min \frac{Cut(C_1, C_2)}{Vol(C_1)} + \frac{Cut(C_1, C_2)}{Vol(C_2)} = \min \left(\frac{1}{Vol(C_1)} + \frac{1}{Vol(C_2)} \right) Cut(C_1, C_2)$$



- Unnormalized cuts are attracted to outliers.

Graph Cut Criteria

- Criterion: **Normalized-cut** [Shi-Malik, '97]
 - Connectivity between groups relative to the density of each group

$$ncut(A, B) = \frac{cut(A, B)}{vol(A)} + \frac{cut(A, B)}{vol(B)}$$

$Vol(A)$: total weight of the edges with at least one endpoint in :

■ Why use this criterion?

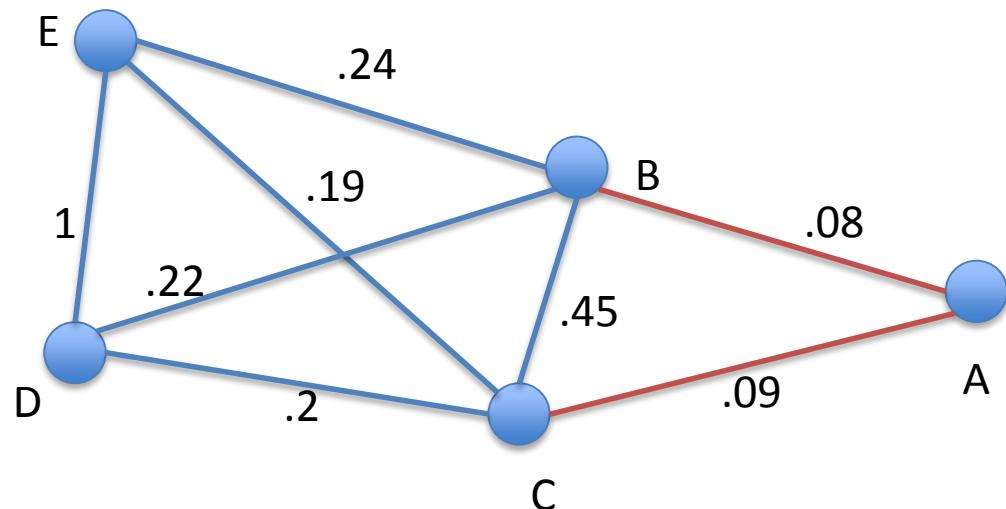
- Produces more balanced partitions
- How do we efficiently find a good partition?
 - Problem: Computing optimal cut is NP-hard

$vol(A)$: total weighted degree of the nodes in A :

$$vol(A) = \sum_{i \in A} k_i \text{ (number of edge end points in } A\text{)}$$

Spectral Clustering Example

- Minimum Cut



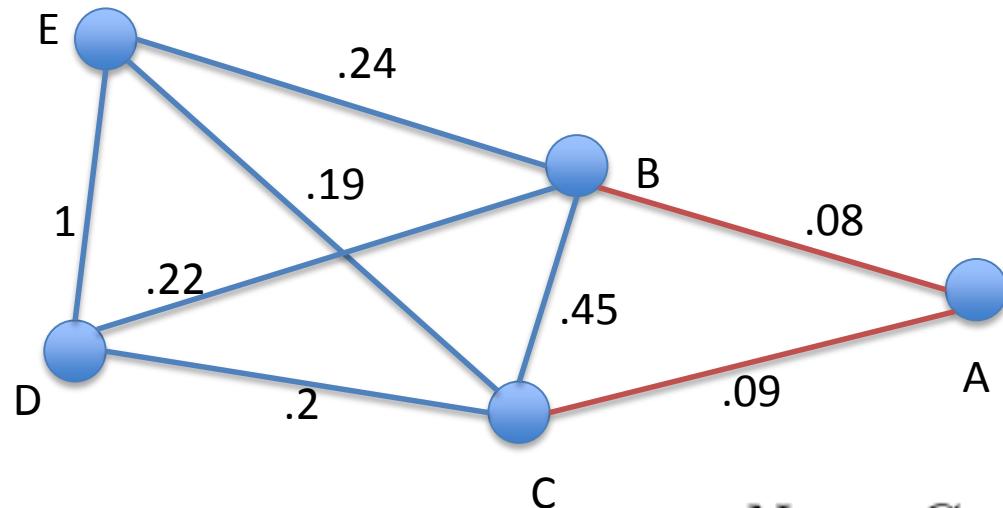
Height	Weight
20	5
8	6
9	4
4	4
4	5

$$Cut(BCDE, A) = 0.17$$

Spectral Clustering Example

- Normalized Minimum Cut

$$NormCut(C_1, C_2) = \frac{Cut(C_1, C_2)}{Vol(C_1)} + \frac{Cut(C_1, C_2)}{Vol(C_2)}$$



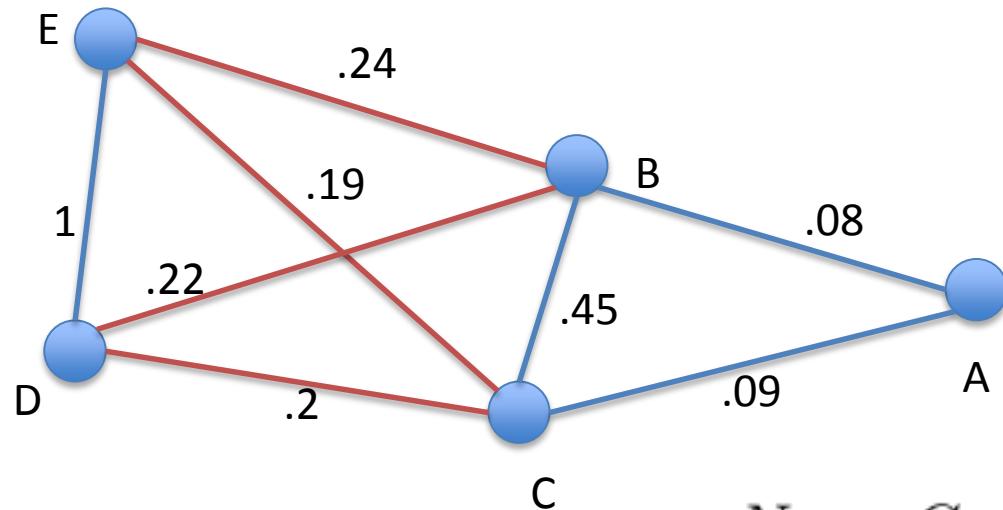
Height	Weight
20	5
8	6
9	4
4	4
4	5

$$NormCut(BCDE, A) = 1.067$$

Spectral Clustering Example

- Normalized Minimum Cut

$$NormCut(C_1, C_2) = \frac{Cut(C_1, C_2)}{Vol(C_1)} + \frac{Cut(C_1, C_2)}{Vol(C_2)}$$



Height	Weight
20	5
8	6
9	4
4	4
4	5

$$NormCut(BCDE, A) = 1.067$$

$$NormCut(ABC, DE) = 1.038$$

Problem

- Identifying a minimum cut is NP-hard.
- There are efficient approximations using linear algebra.
- Based on the Laplacian Matrix, or **graph Laplacian**

Spectral Graph Partitioning

- A : adjacency matrix of undirected \mathbf{G}
 - $A_{ij} = 1$ if i is an edge, else 0
- x is a vector in \mathbb{R}^n with components
 - Think of it as a label/value of each node of
- **What is the meaning of $A \cdot x$?**

$$\begin{bmatrix} a_{11} & \dots & a_{1n} \\ \vdots & & \vdots \\ a_{n1} & \dots & a_{nn} \end{bmatrix} \begin{bmatrix} x_1 \\ \vdots \\ x_n \end{bmatrix} = \begin{bmatrix} y_1 \\ \vdots \\ y_n \end{bmatrix}$$

$$y_i = \sum_{j=1}^n A_{ij} x_j = \sum_{(i,j) \in E} x_j$$

- **Entry y_i is a sum of labels x_j of neighbors of i**

What is the meaning of Ax ?

- **j^{th} coordinate of $A \cdot x$:**

- Sum of the x -values
of neighbors of j

- Make this a new value at node j

$$\begin{bmatrix} a_{11} & \dots & a_{1n} \\ \vdots & & \vdots \\ a_{n1} & \dots & a_{nn} \end{bmatrix} \begin{bmatrix} x_1 \\ \vdots \\ x_n \end{bmatrix} = \lambda \begin{bmatrix} x_1 \\ \vdots \\ x_n \end{bmatrix}$$

- **Spectral Graph Theory:**

- Analyze the “spectrum” of matrix representing

- **Spectrum:** Eigenvectors of a graph, ordered by
the magnitude (strength) of their corresponding
eigenvalues λ_i : $\Lambda = \{\lambda_1, \lambda_2, \dots, \lambda_n\}$

$$\lambda_1 \leq \lambda_2 \leq \dots \leq \lambda_n$$

Note : we sort Lambda's in ascending (Not Descending) Order

Example | d - Regular Graph

- Suppose all nodes in G have degree d (G is d -regular) and G is connected
- What are some eigenvalues/vectors of G ?

$$A \cdot x = \lambda \cdot x \quad \text{What is } \lambda? \text{ What } x?$$

- Let's try: $x = (1, 1, \dots, 1)$
- Then: $A \cdot x = (d, d, \dots, d) = \lambda \cdot x$. So: $\lambda = d$
- We found an eigenpair of G :
 $x = (1, 1, \dots, 1), \lambda = d$
- d is the largest eigenvalue of A (see next slide)

Note, this is just one eigenpair.

An n by n matrix can have up to n eigenpairs.

Remember the meaning of $y = A \cdot x$:

$$y_i = \sum_{j=1}^n A_{ij}x_j = \sum_{(i,j) \in E} x_j$$

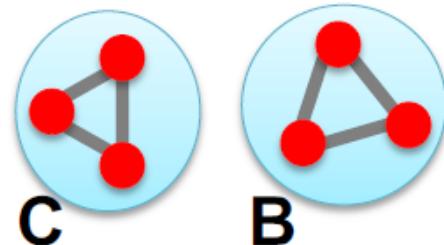
d is the largest eigen value of A

- G is d -regular connected, A is its adjacency matrix
- **Claim:**
 - (1) d has multiplicity of 1 (there is only 1 eigenvector associated with eigenvalue d)
 - (2) d is the largest eigenvalue of A
- **Proof:**
 - To obtain value eigval d we needed $x_i = x_j$ for every i, j
 - This means $x = c \cdot (1, 1, \dots, 1)$ for some const. c
 - **Define:** Set $S = \text{nodes } i \text{ with maximum value of } x_i$
 - Then consider some vector y which is not a multiple of vector $(1, \dots, 1)$. So not all nodes i (with labels y_i) are in S
 - Consider some node $j \in S$ and a neighbor $i \notin S$ then node j gets a value strictly less than d
 - So y is not eigenvector! And so d is the largest eigenvalue!

Example: Graph on 2 Components

- **What if G is not connected?**

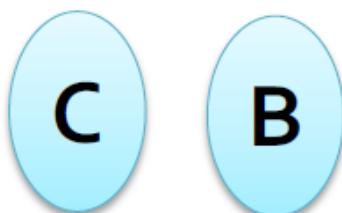
- G has 2 components, each d -regular



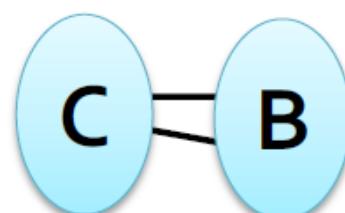
- **What are some eigenvectors?**

- $x = \text{Put all } 1\text{s on } C \text{ and } 0\text{s on } B \text{ or vice versa}$
 - $x' = (\underbrace{1, \dots, 1}_{|C|}, \underbrace{0, \dots, 0}_{|B|})^T$ then $A \cdot x' = (d, \dots, d, 0, \dots, 0)^T$
 - $x'' = (0, \dots, 0, 1, \dots, 1)^T$ then $A \cdot x'' = (0, \dots, 0, d, \dots, d)^T$
 - And so in both cases the corresponding $\lambda = d$

- **A bit of intuition:**



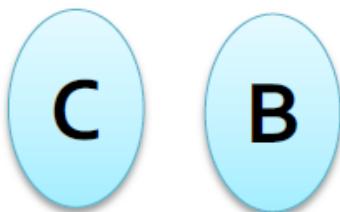
$$\lambda_n = \lambda_{n-1}$$



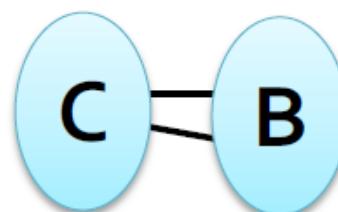
$$\lambda_n - \lambda_{n-1} \approx 0$$

2nd largest eigenval.
 λ_{n-1} now has
value very close
to λ_n

More intuition



$$\lambda_n = \lambda_{n-1}$$



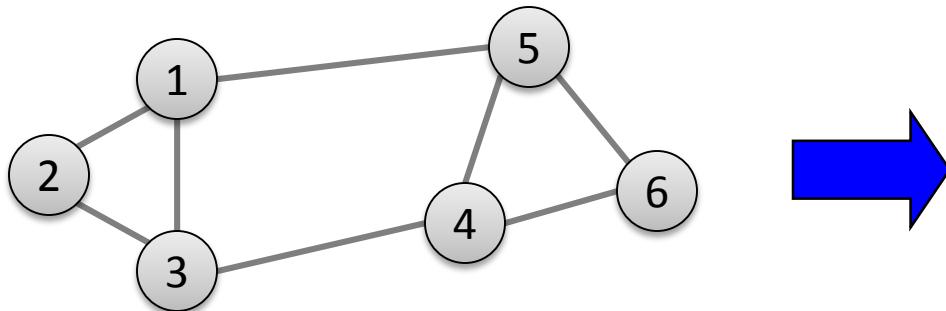
$$\lambda_n - \lambda_{n-1} \approx 0$$

2nd largest eigval.
 λ_{n-1} now has
value very close
to λ_n

- If the d-regular graph is connected (right example) then we already know that $x_n = (1, \dots, 1)$ is an eigenvector
- Eigenvectors are orthogonal so then the components of x_{n-1} must sum to 0
 - Why? $x_n \cdot x_{n-1} = 0$ then $\sum_i x_n[i] \cdot x_{n-1}[i] = \sum_i x_{n-1}[i] = 0$
 - x_{n-1} “splits” the nodes into two groups
 - $x_{n-1}[i] > 0$ vs. $x_{n-1}[i] < 0$
 - So we in principle could look at the eigenvector of the 2nd largest eigenvalue and declare nodes with positive label in **C** and negative label in **B**. (but there are still many details for us to figure out here)

Matrix Representations

- **Adjacency matrix (A):**
 - $n \times n$ matrix
 - $A = [a_{ij}]$, $a_{ij} = 1$ if edge between node i and j

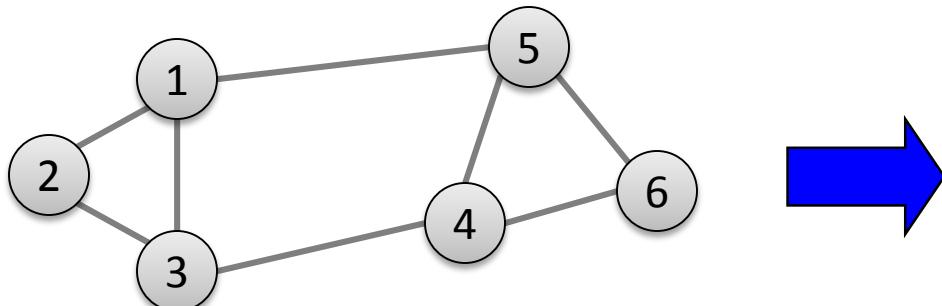


	1	2	3	4	5	6
1	0	1	1	0	1	0
2	1	0	1	0	0	0
3	1	1	0	1	0	0
4	0	0	1	0	1	1
5	1	0	0	1	0	1
6	0	0	0	1	1	0

- **Important properties:**
 - Symmetric matrix
 - Eigenvectors are real and orthogonal

Matrix Representations

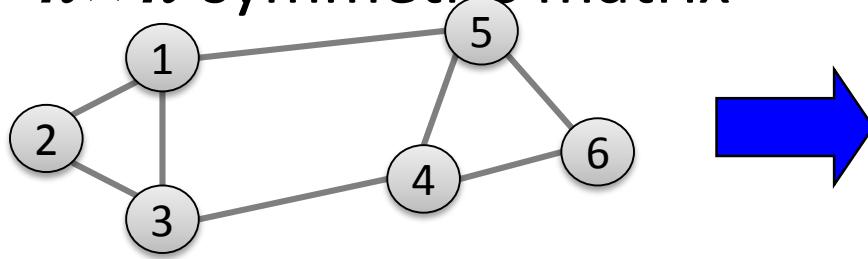
- **Degree matrix (D):**
 - $n \times n$ diagonal matrix
 - $D = [d_{ii}]$, d_{ii} = degree of node i



	1	2	3	4	5	6
1	3	0	0	0	0	0
2	0	2	0	0	0	0
3	0	0	3	0	0	0
4	0	0	0	3	0	0
5	0	0	0	0	3	0
6	0	0	0	0	0	2

Matrix Representations

- **Laplacian matrix (L):**
 - $n \times n$ symmetric matrix



	1	2	3	4	5	6
1	3	-1	-1	0	-1	0
2	-1	2	-1	0	0	0
3	-1	-1	3	-1	0	0
4	0	0	-1	3	-1	-1
5	-1	0	0	-1	3	-1
6	0	0	0	-1	-1	2

$$L = D - A$$

- **What is trivial eigenpair?**
- **Important properties:**
 - **Eigenvalues** are non-negative real numbers
 - **Eigenvectors** are real and orthogonal

3 Facts about The Laplacian L

(a) All eigenvalues are ≥ 0

(b) $x^T L x = \sum_{ij} L_{ij} x_i x_j \geq 0$ for every x

(c) L can be written as $L = N^T \cdot N$

- That is, L is positive semi-definite

■ **Proof: (the 3 facts are saying the same thing)**

- (c) \Rightarrow (b): $x^T L x = x^T N^T N x = (Nx)^T (Nx) \geq 0$

- As it is just the square of length of Nx

- (b) \Rightarrow (a): Let λ be an eigenvalue of L . Then by (b)
 $x^T L x \geq 0$ so $x^T L x = x^T \lambda x = \lambda x^T x \Rightarrow \lambda \geq 0$

- (a) \Rightarrow (c): is also easy! Do it yourself.

L2 as Optimization Problem

- Fact: For symmetric matrix M :

$$\lambda_2 = \min_{x : x^T w_1 = 0} \frac{x^T M x}{x^T x}$$

See next slide
for the proof.
Deriving this
is a HW problem

(w_1 is eigenvector corresponding to λ_1)

- What is the meaning of $\min x^T L x$ on G ?

$$\begin{aligned} x^T L x &= \sum_{i,j=1}^n L_{ij} x_i x_j = \sum_{i,j=1}^n (D_{ij} - A_{ij}) x_i x_j \\ &= \sum_i D_{ii} x_i^2 - \sum_{(i,j) \in E} 2x_i x_j \\ &= \sum_{(i,j) \in E} (\underbrace{x_i^2 + x_j^2}_{\text{Node } i \text{ has degree } d_i} - 2x_i x_j) = \sum_{(i,j) \in E} (\mathbf{x}_i - \mathbf{x}_j)^2 \end{aligned}$$

Node i has degree d_i . So, value x_i^2 needs to be summed up d_i times.
But each edge (i,j) has two endpoints so we need $x_i^2 + x_j^2$

Proof: $\lambda_2 = \min_{x : x^T w_1 = 0} \frac{x^T M x}{x^T x}$

- Write x in basis of eigenvecs w_1, w_2, \dots, w_n of M and λ_i are corresponding eigenvalues. So, $x = \sum_i^n \alpha_i w_i$
- Then we get: $Mx = \sum_i \alpha_i M w_i = \sum_i \alpha_i \lambda_i w_i$
- So, what is $x^T M x$?

$$\begin{aligned} x^T M x &= (\underbrace{\sum_i \alpha_i w_i}_x)^T (\underbrace{\sum_i \alpha_i \lambda_i w_i}_M) = \sum_{ij} \alpha_i \lambda_j \alpha_j \\ &= \sum_i \alpha_i^2 \lambda_i \underbrace{w_i^T w_i}_{= 0 \text{ if } i \neq j, 1 \text{ otherwise}} = \sum_i \lambda_i \alpha_i^2 \end{aligned}$$

- Want minimize this over all unit vectors w :
- $w = \min$ over choices of $(\alpha_1, \dots, \alpha_n)$ so that:
 - $x^T w_1 = 0$, rewrite it as $(\sum_i \alpha_i w_i) \cdot w_1 = 0$ and remember that $w_i^T w_j = 0$ (because w are eigenvectors). Then $\alpha_1 = 0$
 - $\sum \alpha_i^2 = 1$ (unit length)
- So, to minimize this, set $\alpha_2 = 1$ and the rest to 0 $\sum_i \lambda_i \alpha_i^2 = \lambda_2$

Finding X that Solves $\lambda_2 = \min_{x : x^T w_1=0} \frac{x^T M x}{x^T x}$

- What else do we know about x ?

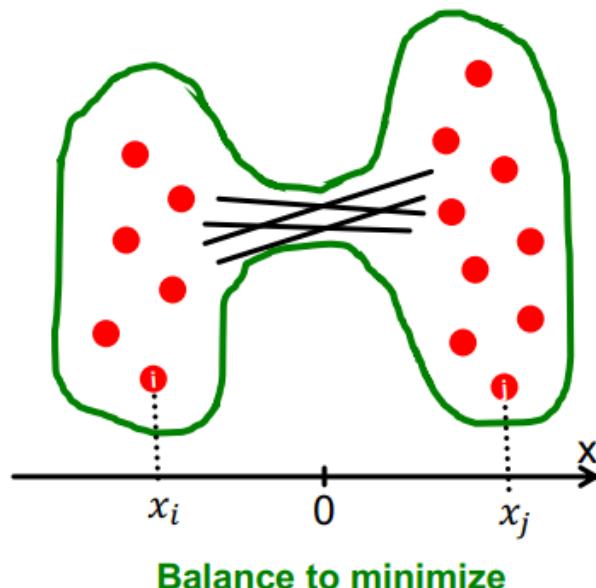
- x is unit vector: $\sum_i x_i^2 = 1$
- x is orthogonal to 1st eigenvector $(1, \dots, 1)$ thus:
 $\sum_i x_i \cdot 1 = \sum_i x_i = 0$

- Remember:

$$\lambda_2 = \min \frac{\sum_{(i,j) \in E} (x_i - x_j)^2}{\sum_i x_i^2}$$

All labelings
 of nodes i so
 that $\sum x_i = 0$

We want to assign values x_i to nodes i such
 that few edges cross 0.
 (we want x_i and x_j to subtract each other)



Finding Optimal Cut (Fidler 73)

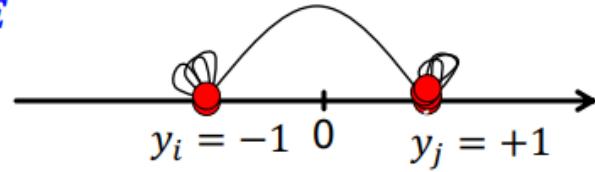
- Back to finding the optimal cut
- Express partition (A,B) as a vector

$$y_i = \begin{cases} +1 & \text{if } i \in A \\ -1 & \text{if } i \in B \end{cases}$$

- Enforce that $|A| = |B| \rightarrow \sum_i y_i = 0$
 - Equivalent to being orthogonal to the trivial eigenvector $(1, \dots, 1)$
- We can minimize the cut of the partition by finding a vector y that minimizes:

$$\arg \min_{y \in \{-1, +1\}^n} f(y) = \sum_{(i,j) \in E} (y_i - y_j)^2$$

Can't solve exactly. Let's relax y and allow it to take any real value.

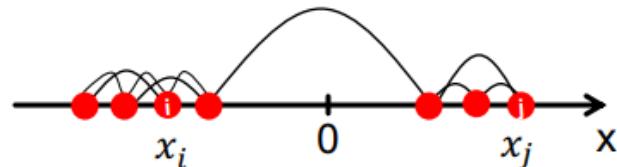


Rayleigh Theorem

$$\lambda_2 = \min_{x : x^T w_1 = 0} \frac{x^T M x}{x^T x}$$

$$\min_{y \in \mathbb{R}^n : \sum_i y_i = 0} f(y) = \sum_{(i,j) \in E} (y_i - y_j)^2 = y^T L y$$

$$\sum_i y_i^2 = 1$$



- $\lambda_2 = \min_y f(y)$: The minimum value of $f(y)$ is given by the 2nd smallest eigenvalue λ_2 of the Laplacian matrix L
- $x = \arg \min_y f(y)$: The optimal solution for y is given by the eigenvector x corresponding to λ_2 , referred to as the **Fiedler vector**
- Can use sign of x_i to determine cluster assignment of node i

Approx. Guarantee of Spectral

- Suppose there is a partition of \mathbf{G} into \mathbf{A} and \mathbf{B} where $|A| \leq |B|$, s.t. “conductance” of the cut (A, B) is $\beta = \frac{(\# \text{ edges from } A \text{ to } B)}{|A|}$ then $\lambda_2 \leq 2\beta$
Note: $|A| < |B|$
 - This is the approximation guarantee of the spectral clustering: Spectral finds a cut that has at most **twice the conductance** as the optimal one of conductance β .
- Proof:**
 - Let: $a = |A|$, $b = |B|$ and $e = \# \text{ edges from } A \text{ to } B$
 - Enough to choose some x_i based on A and B such that:
$$\lambda_2 \leq \underbrace{\frac{\sum(x_i - x_j)^2}{\sum_i x_i^2}}_{\leq 2\beta} \leq 2\beta$$
 (while also $\sum_i x_i = 0$)

Approx. Guarantee of Spectral

■ Proof (continued):

■ 1) Let's set: $x_i = \begin{cases} -\frac{1}{a} & \text{if } i \in A \\ +\frac{1}{b} & \text{if } i \in B \end{cases}$

Note: $|A| < |B|$

■ Let's quickly verify that $\sum_i x_i = 0$: $a\left(-\frac{1}{a}\right) + b\left(\frac{1}{b}\right) = 0$

■ 2) Then: $\frac{\sum(x_i - x_j)^2}{\sum_i x_i^2} = \frac{\sum_{i \in A, j \in B} \left(\frac{1}{b} + \frac{1}{a}\right)^2}{a\left(-\frac{1}{a}\right)^2 + b\left(\frac{1}{b}\right)^2} = \frac{e \cdot \left(\frac{1}{a} + \frac{1}{b}\right)^2}{\frac{1}{a} + \frac{1}{b}} =$

$e \left(\frac{1}{a} + \frac{1}{b}\right) \leq e \left(\frac{1}{a} + \frac{1}{a}\right) = e \frac{2}{a} \leq 2\beta$

e ... number of edges between A and B

Which proves that the cost achieved by spectral is better than twice the OPT cost

Approx. Guarantee of Spectral

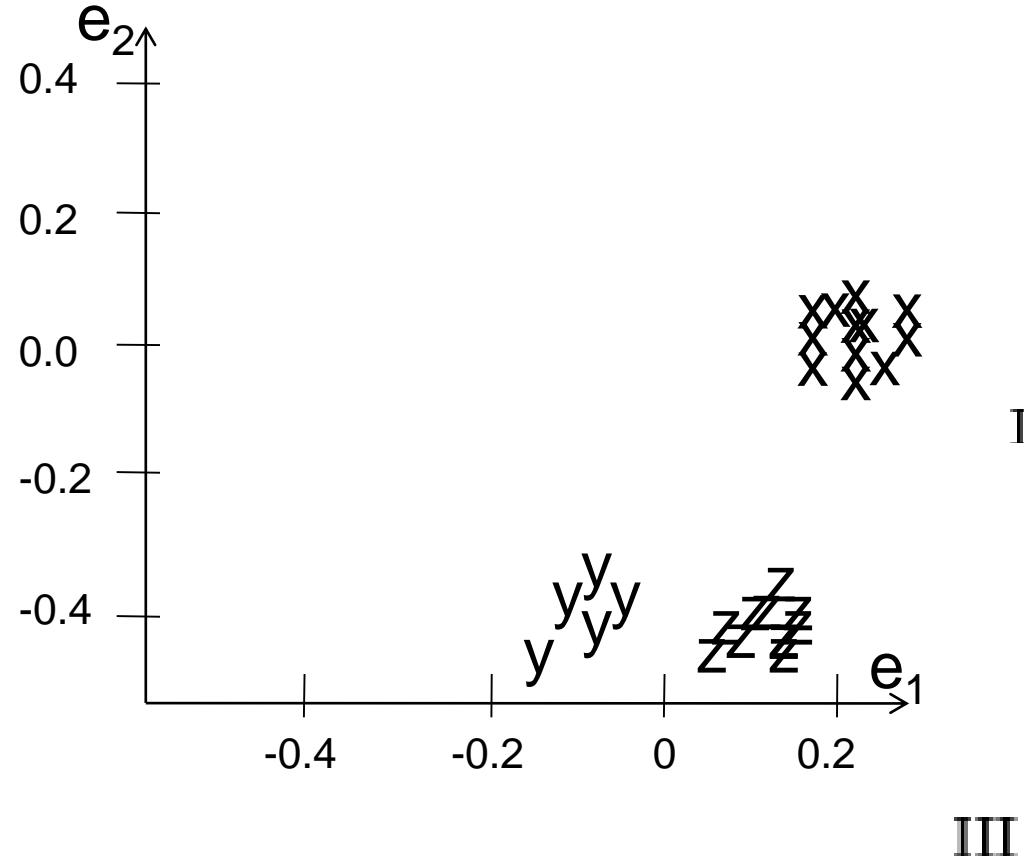
- Putting it all together: The Cheeger inequality

$$\frac{\beta^2}{2k_{max}} \leq \lambda_2 \leq 2\beta$$

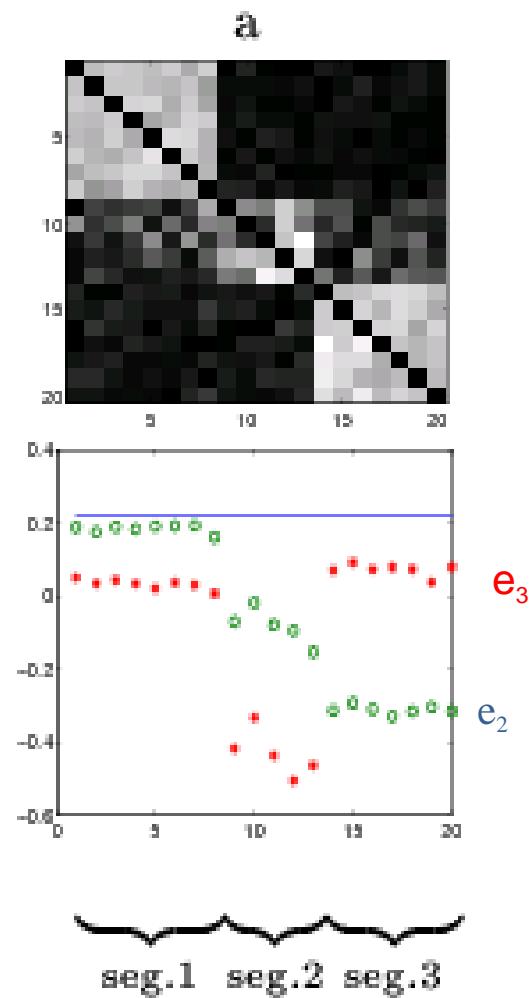
- where k_{max} is the maximum node degree in the graph
 - Note we only provide the 1st part: $\lambda_2 \leq 2\beta$
 - We did not prove $\frac{\beta^2}{2k_{max}} \leq \lambda_2$
- Overall this always certifies that λ_2 always gives a useful bound

Spectral Clustering: Graph = Matrix

$W^*v_1 = v_2$ “propogates weights from neighbors”



[Shi & Meila, 2002]



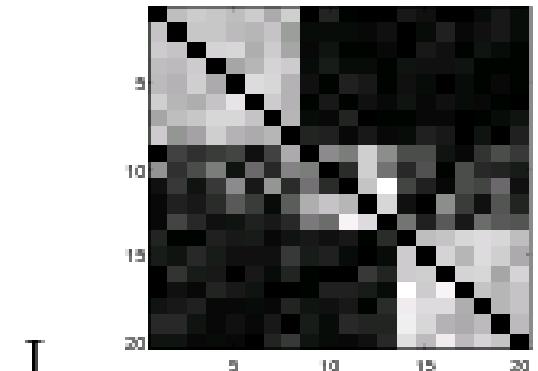
Spectral Clustering: Graph = Matrix

$W^*v_1 = v_2$ “propagates weights from neighbors”

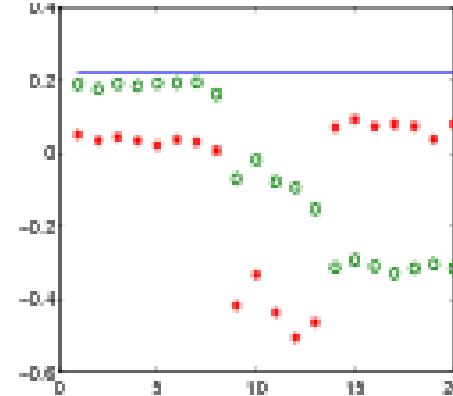
$\mathbf{W} \cdot \mathbf{v} = \lambda \mathbf{v}$: \mathbf{v} is an eigenvector with eigenvalue λ

If W is connected but roughly block diagonal with k blocks then

- the top eigenvector is a constant vector
- the next k eigenvectors are roughly piecewise constant with “pieces” corresponding to blocks



I



III



M

Spectral Clustering: Graph = Matrix

$\mathbf{W}^* \mathbf{v}_1 = \mathbf{v}_2$ “propagates weights from neighbors”

$\mathbf{W} \cdot \mathbf{v} = \lambda \mathbf{v}$: \mathbf{v} is an eigenvector with eigenvalue λ

If \mathbf{W} is connected but roughly block diagonal with k blocks then

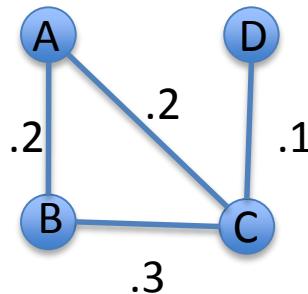
- the “top” eigenvector is a constant vector
- the next k eigenvectors are roughly piecewise constant with “pieces” corresponding to blocks

Spectral clustering:

- Find the top $k+1$ eigenvectors $\mathbf{v}_1, \dots, \mathbf{v}_{k+1}$
- Discard the “top” one
- Replace every node a with k -dimensional vector $x_a = \langle \mathbf{v}_2(a), \dots, \mathbf{v}_{k+1}(a) \rangle$
- Cluster with k -means

Spectral Clustering

- Construct an **affinity** matrix



$W =$

	A	B	C	D
A	.4	.2	.2	0
B	.2	.5	.3	0
C	.2	.3	.6	.1
D	0	0	.1	.1

Spectral Clustering

- Construct the graph Laplacian

$$D = \text{diag}(D_0, D_1, \dots, D_n)$$

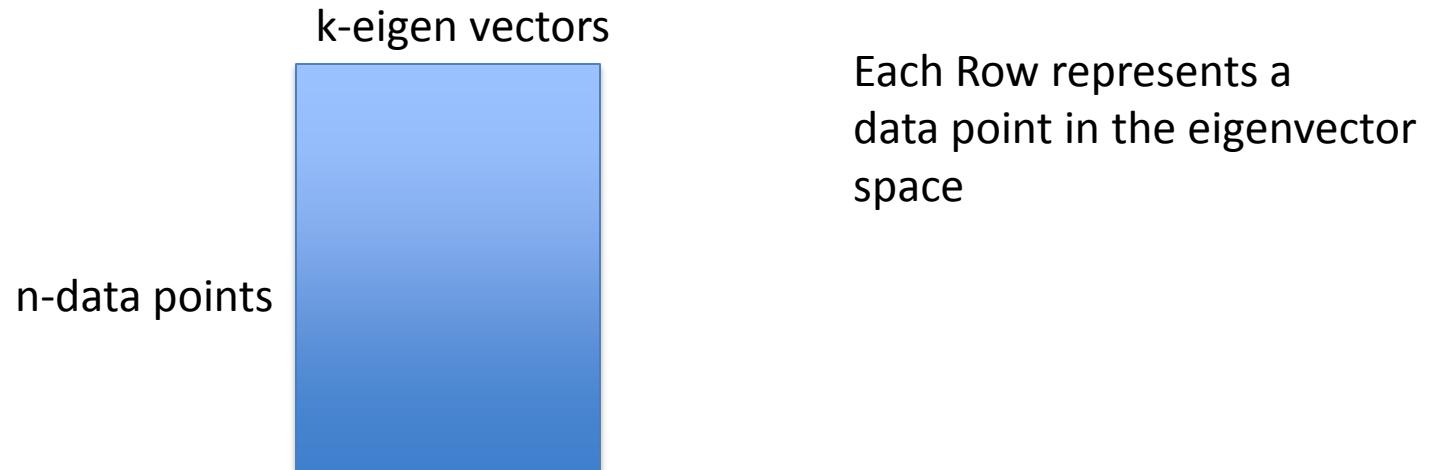
$$L = D - W$$

- Identify eigenvectors of the affinity matrix

$$Lv = \lambda v$$

Spectral Clustering

- K-Means on eigenvector transformation of the data.



- Project back to the initial data representation.

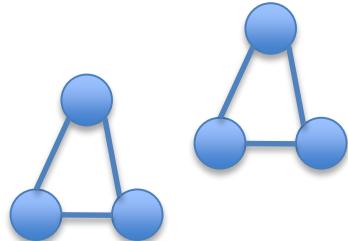
Overview: what are we doing?

- Define the affinity matrix
- Identify eigenvalues and eigenvectors.
- K-means of transformed data
- Project back to original space

Why does this work?

- Ideal Case

$$Lv = \lambda v$$



1	1	1	0	0	0
1	1	1	0	0	0
1	1	1	0	0	0
0	0	0	1	1	1
0	0	0	1	1	1
0	0	0	1	1	1

1	0
1	0
1	0
0	1
0	1
0	1

- What are we optimizing? Why do the eigenvectors of the laplacian include cluster identification information

Normalized Graph Cuts view

- Minimal (bipartitional) normalized cut.

$$\min \frac{Cut(C_1, C_2)}{Vol(C_1)} + \frac{Cut(C_1, C_2)}{Vol(C_2)} = \min \left(\frac{1}{Vol(C_1)} + \frac{1}{Vol(C_2)} \right) Cut(C_1, C_2)$$

$$NCut(A, B) = \frac{y^T(D - W)y}{y^T D y}$$

$$\min_y y^T(D - W)y \text{ subject to } y^T D y = 1$$

$$(D - W)y = \lambda D y$$

- Eigenvalues of the laplacian are approximate solutions to mincut problem.

The Laplacian Matrix

- $L = D - W$
- Positive semi-definite $x^T L x \geq 0$
- The lowest eigenvalue is 0, eigenvector is $\vec{1}$
- The second lowest contains the solution
 - The corresponding eigenvector contains the cluster indicator for each data point

$$\lambda_2 = \frac{Cut(A, B)}{|A|} + \frac{Cut(A, B)}{|B|}$$

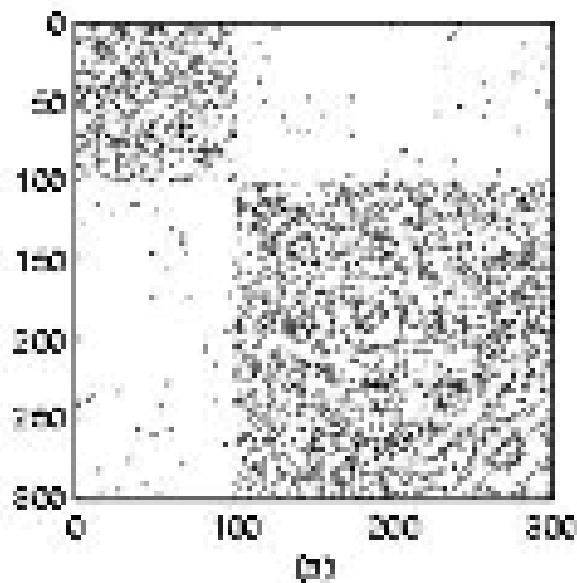
Using eigenvectors to partition

- Each eigenvector partitions the data set into two clusters.
- The entry in the second eigenvector determines the first cut.
- Subsequent eigenvectors can be used to further partition into more sets.

Example

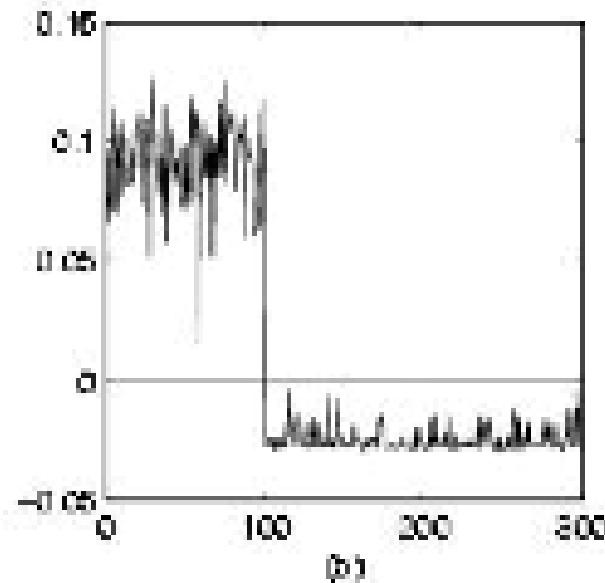
- Dense clusters with some sparse connections

Adjacency matrix



(a)

Eigenvector q_2



(b)

3 class Example

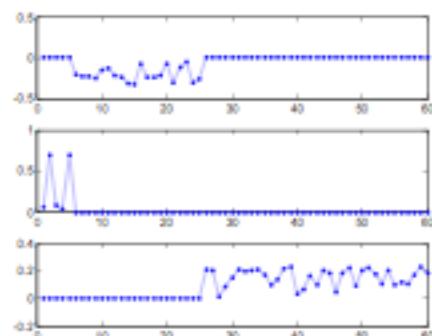
Affinity matrix

$$W; A$$



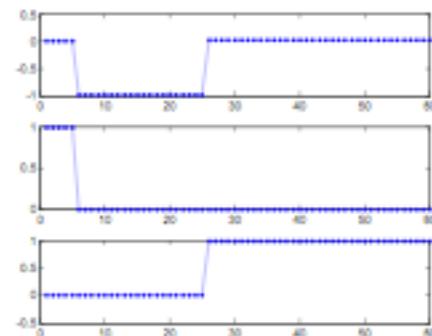
eigenvectors

$$V = [v_1, v_2, v_3]$$

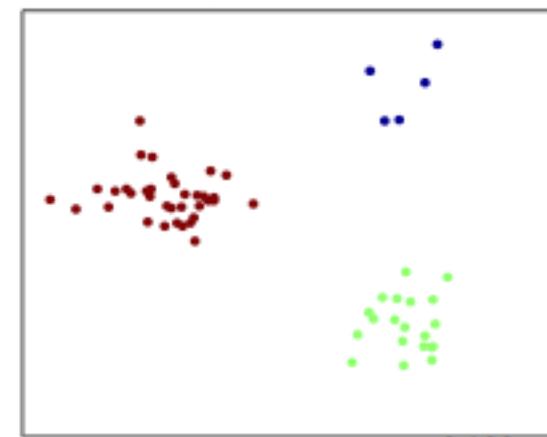


row normalization

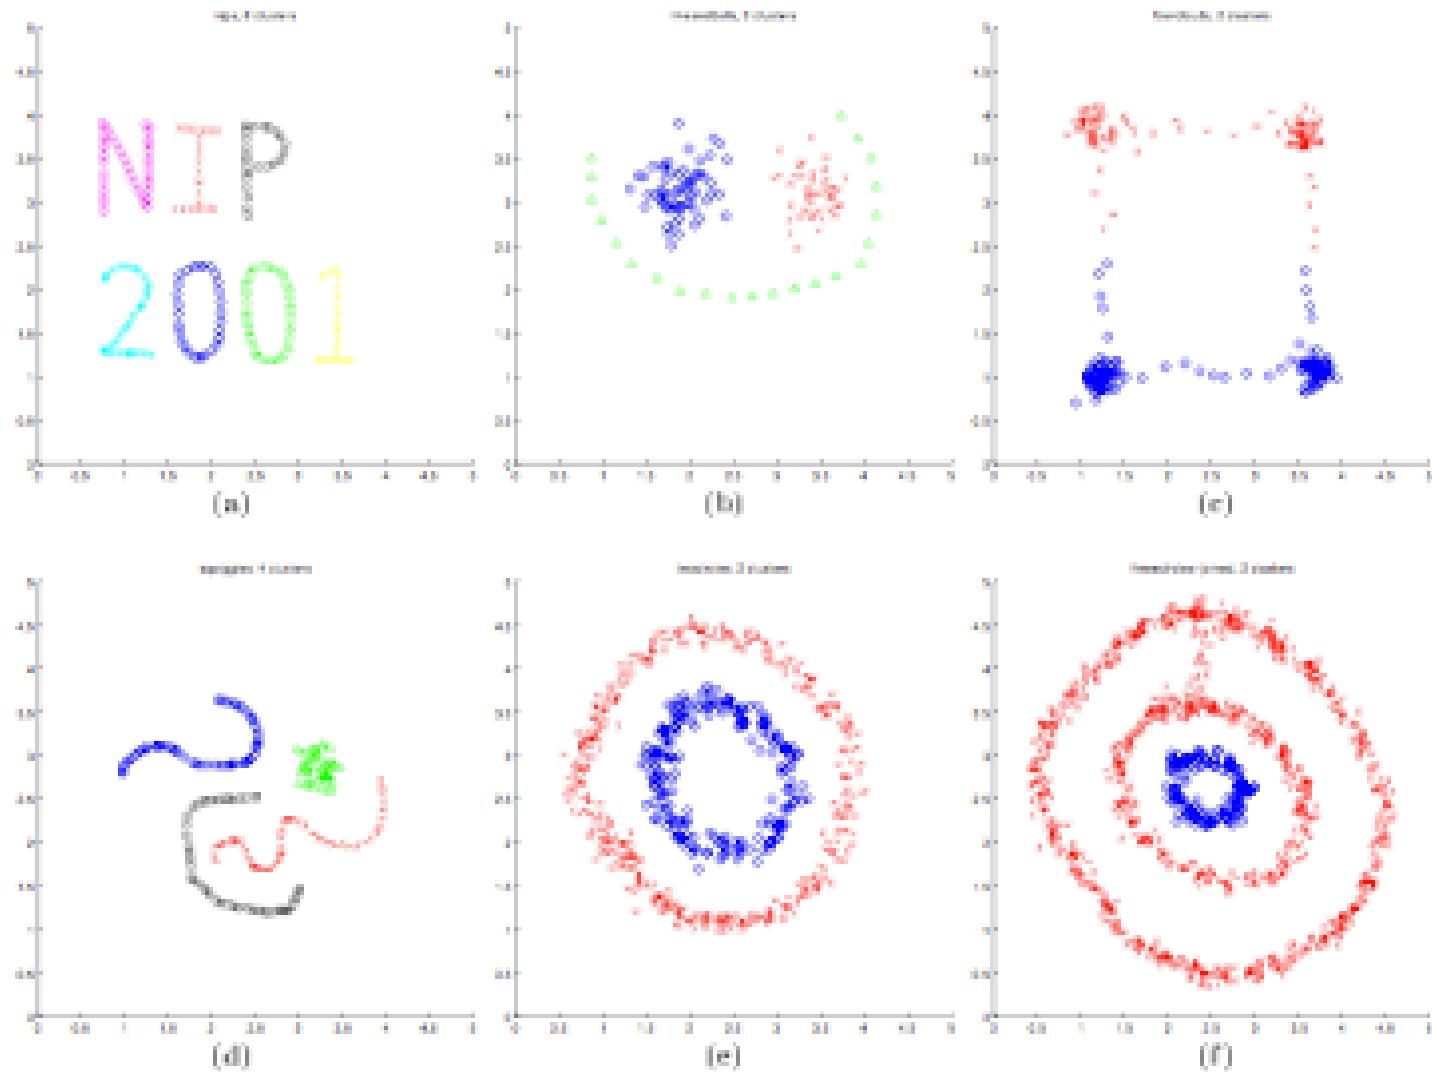
$$U = [u_1, u_2, u_3]$$



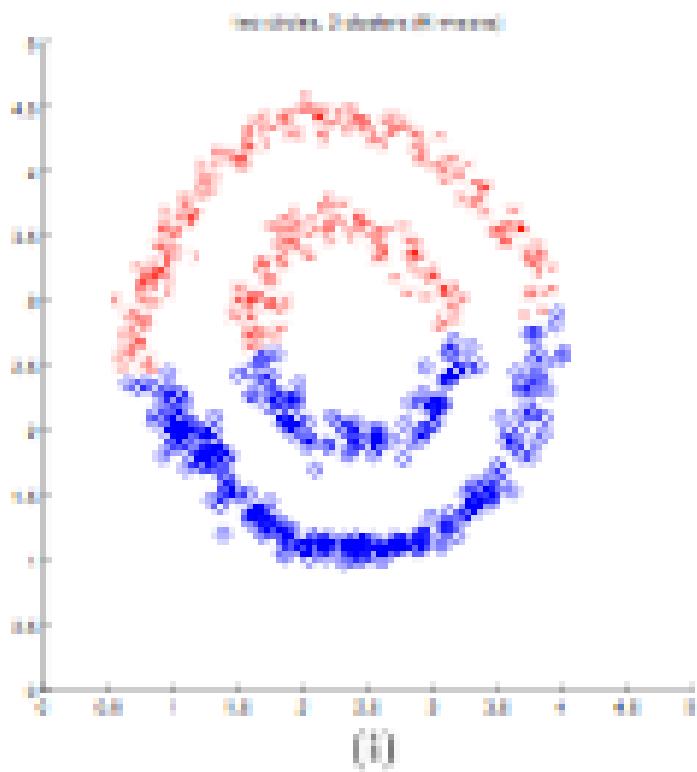
output



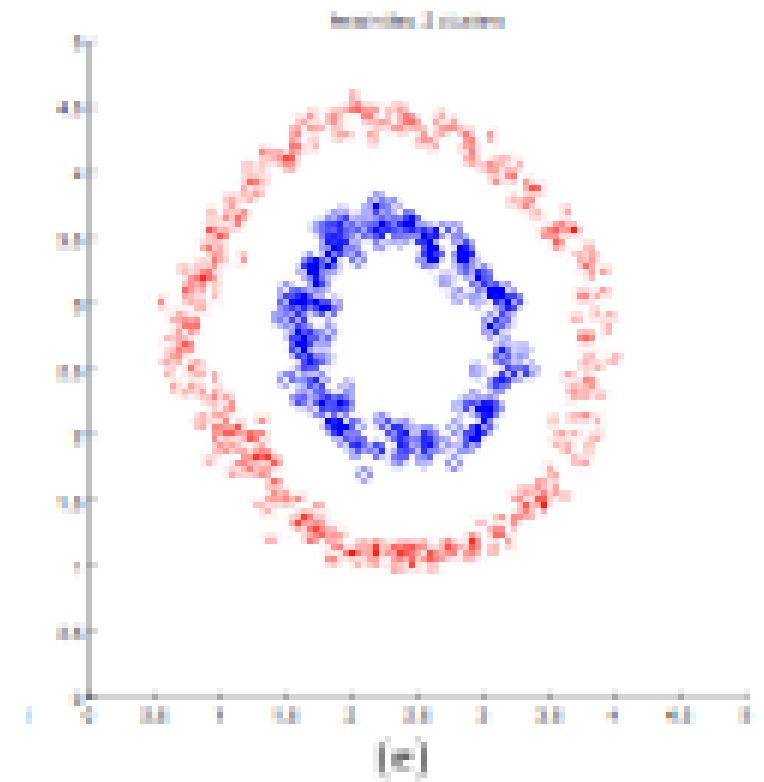
Example [Ng et al. 2001]



k-means vs. Spectral Clustering



K-means



Spectral Clustering

So far...

- **How to define a “good” partition of a graph?**
 - Minimize a given graph cut criterion
- **How to efficiently identify such a partition?**
 - Approximate using information provided by the eigenvalues and eigenvectors of a graph
- **Spectral Clustering**

Spectral Clustering Algorithms

- **Three basic stages:**
 - **1) Pre-processing**
 - Construct a matrix representation of the graph
 - **2) Decomposition**
 - Compute eigenvalues and eigenvectors of the matrix
 - Map each point to a lower-dimensional representation based on one or more eigenvectors
 - **3) Grouping**
 - Assign points to one (or more) clusters, based on the new representation

■ 1) Pre-processing:

- Build Laplacian matrix L of the graph



	1	2	3	4	5	6
1	3	-1	-1	0	-1	0
2	-1	2	-1	0	0	0
3	-1	-1	3	-1	0	0
4	0	0	-1	3	-1	-1
5	-1	0	0	-1	3	-1
6	0	0	0	-1	-1	2

■ 2)

Decomposition:

- Find eigenvalues λ and eigenvectors x of the matrix L



$$\lambda =$$

0.0
1.0
3.0
3.0
4.0
5.0

$$X =$$

0.4	0.3	-0.5	-0.2	-0.4	-0.5
0.4	0.6	0.4	-0.4	0.4	0.0
0.4	0.3	0.1	0.6	-0.4	0.5
0.4	-0.3	0.1	0.6	0.4	-0.5
0.4	-0.3	-0.5	-0.2	0.4	0.5
0.4	-0.6	0.4	-0.4	-0.4	0.0

- Map vertices to corresponding components of x_2

1	0.3
2	0.6
3	0.3
4	-0.3
5	-0.3
6	-0.6



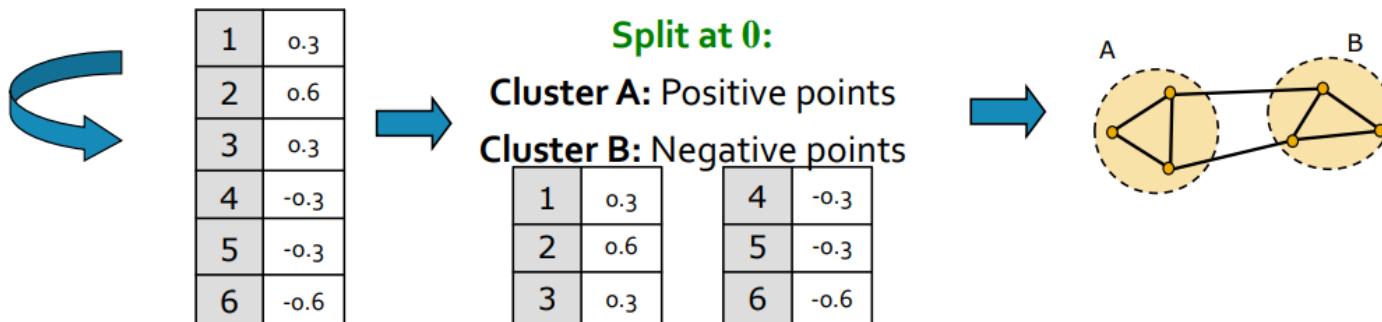
How do we now find the clusters?

■ 3) Grouping:

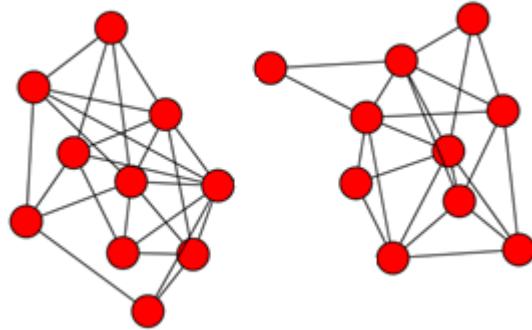
- Sort components of reduced 1-dimensional vector
- Identify clusters by splitting the sorted vector in two

■ How to choose a splitting point?

- Naïve approaches:
 - Split at **0** or median value
- More expensive approaches:
 - Attempt to minimize normalized cut in 1-dimension (sweep over ordering of nodes induced by the eigenvector)



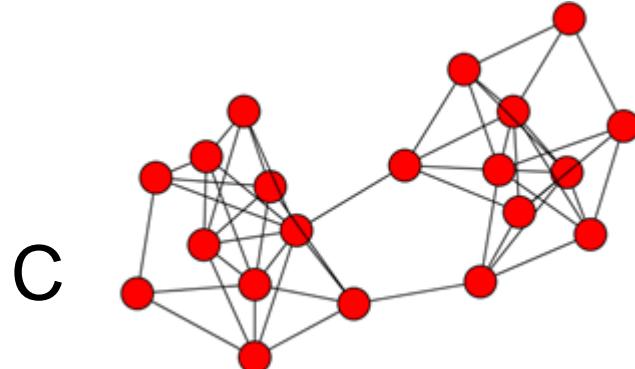
Graphs and Eigenvalues



$$0 = \lambda_1 = \lambda_2$$

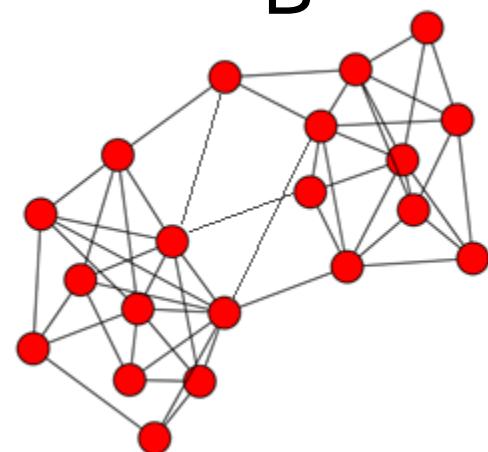
A

$$\lambda_2(C) < \lambda_2(B)$$



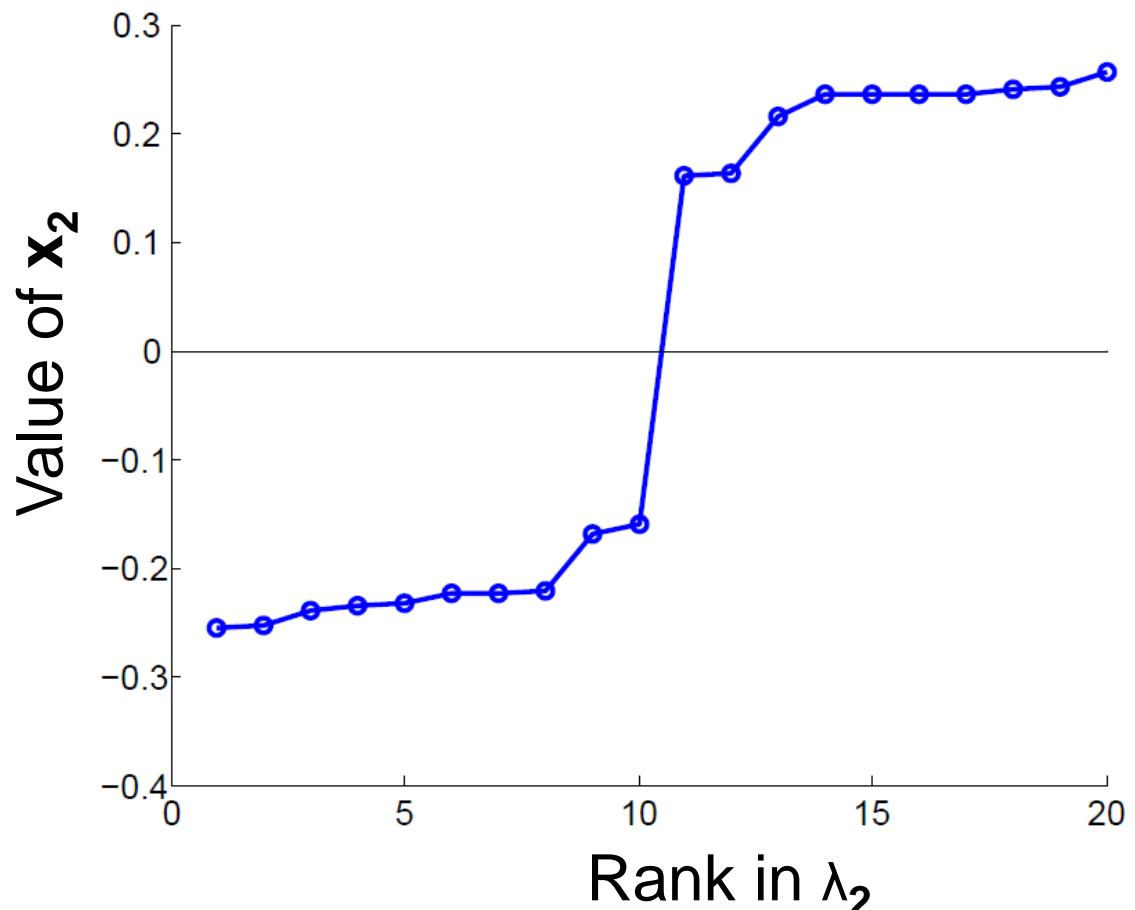
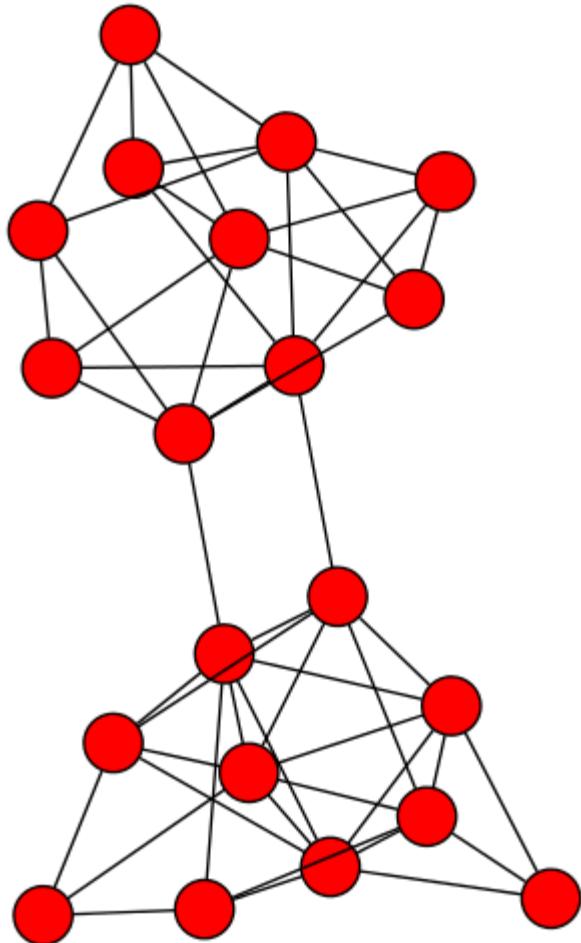
C

$$0 = \lambda_1, \lambda_2 > 0$$

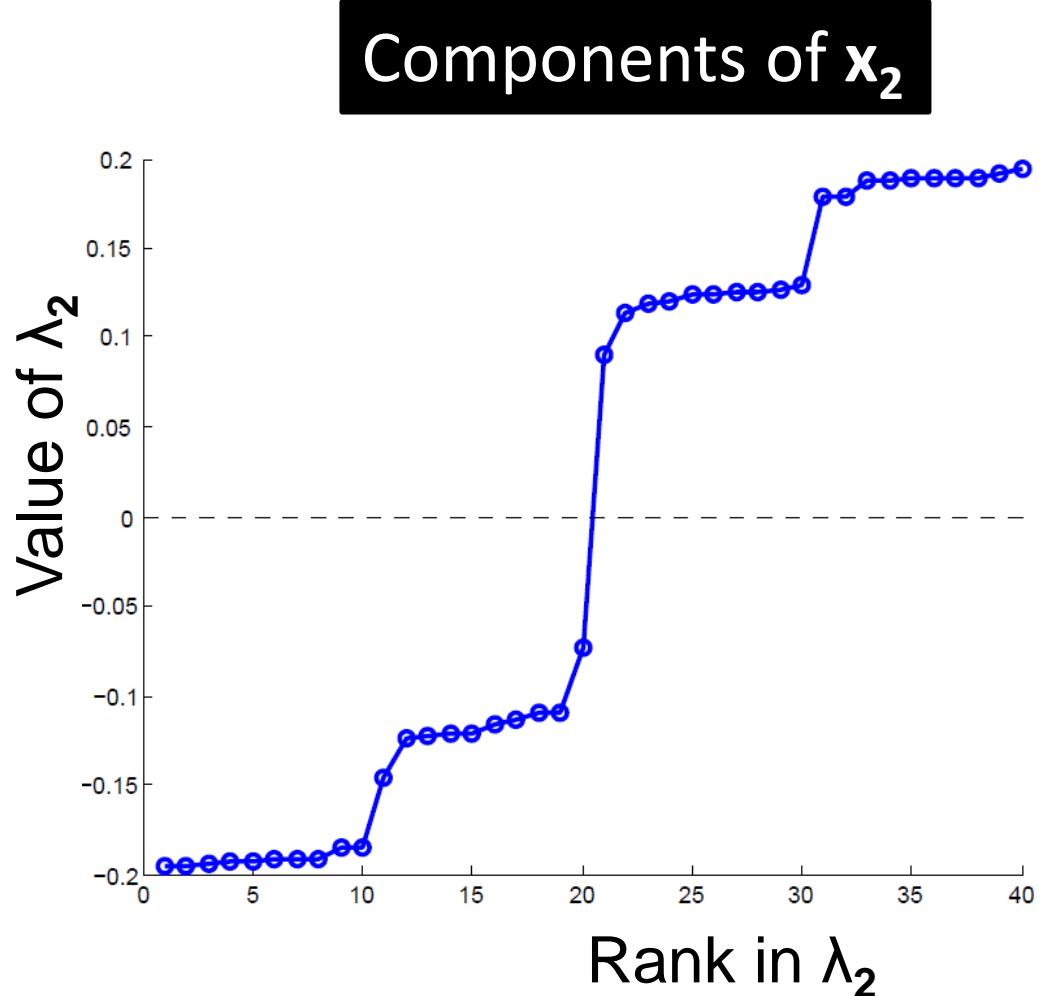
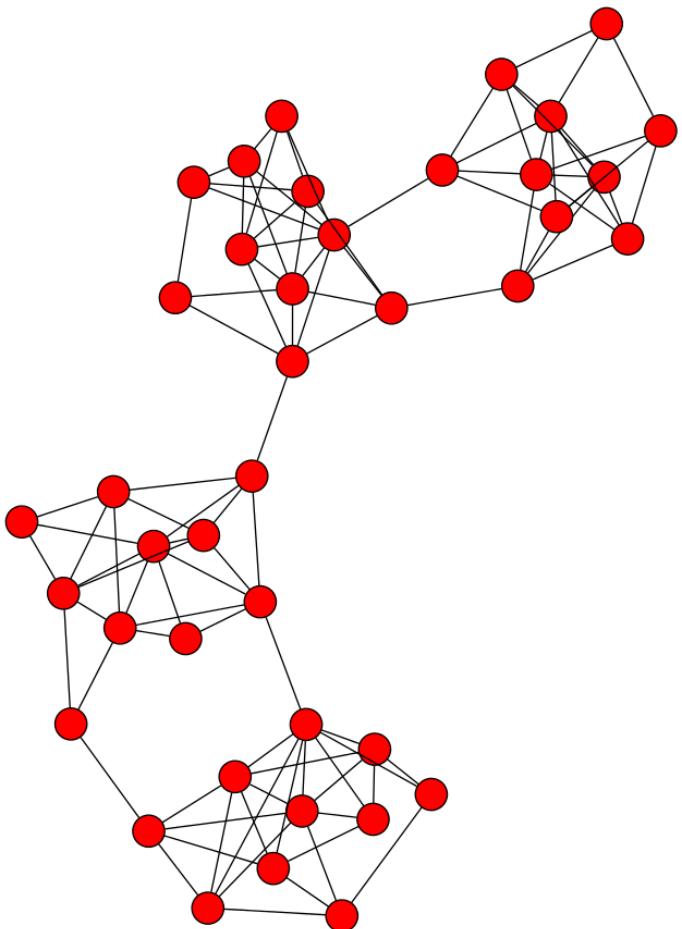


$$0 = \lambda_1, \lambda_2 > 0$$

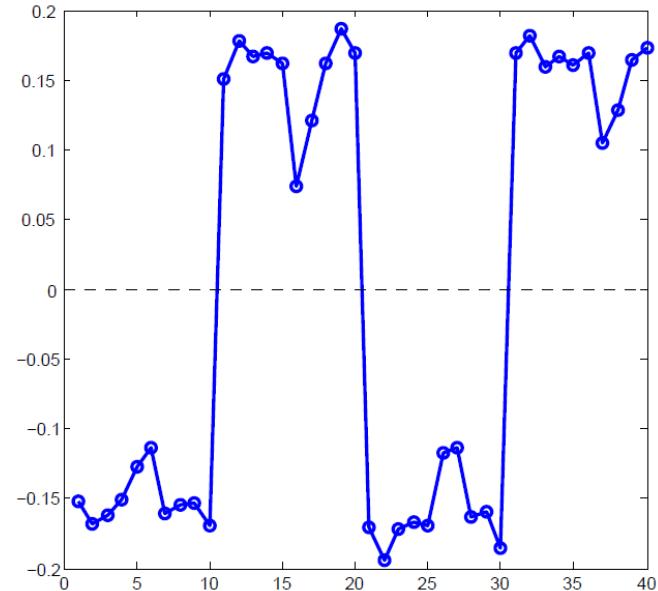
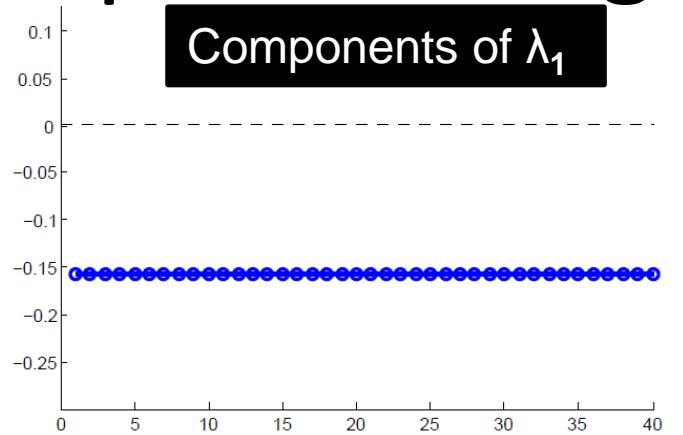
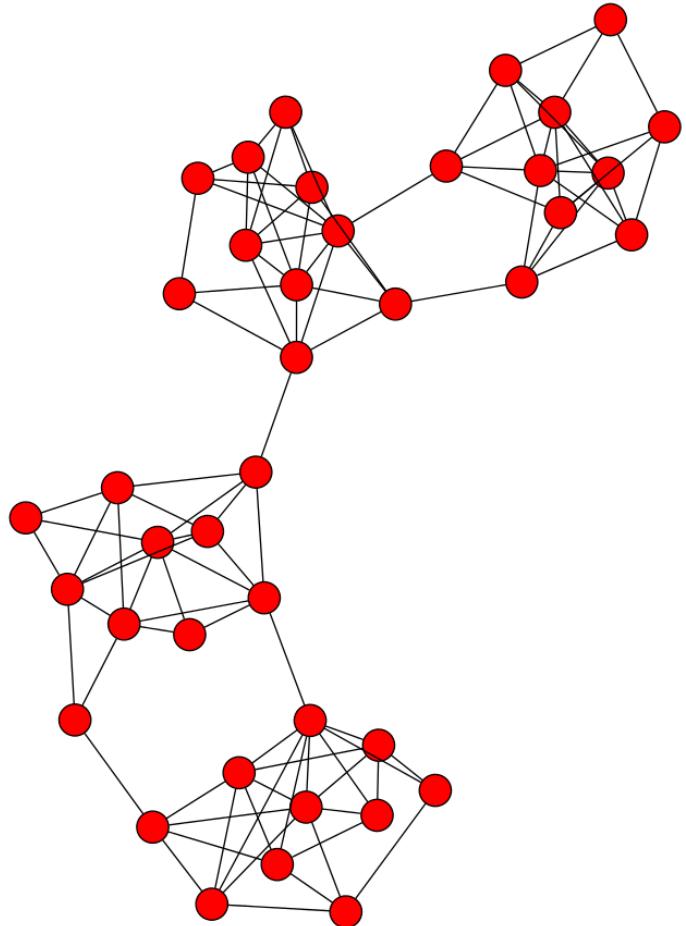
Example: Spectral Partitioning



Example: Spectral Partitioning



Example: Spectral partitioning



■ How do we partition a graph into k clusters?

■ Two basic approaches:

- **Recursive bi-partitioning** [Hagen et al., '92]
 - Recursively apply bi-partitioning algorithm in a hierarchical divisive manner
 - Disadvantages: Inefficient, unstable
- **Cluster multiple eigenvectors** [Shi-Malik, '00]
 - Build a reduced space from multiple eigenvectors
 - Each node is now represented by k numbers
 - We then cluster (apply k-means) the nodes based on their k -dim representation
 - Commonly used in recent papers
 - A preferable approach...

k-Way Spectral Clustering

- **How do we partition a graph into k clusters?**
- **Two basic approaches:**
 - **Recursive bi-partitioning** [Hagen et al., '92]
 - Recursively apply bi-partitioning algorithm in a hierarchical divisive manner
 - Disadvantages: Inefficient, unstable
 - **Cluster multiple eigenvectors** [Shi-Malik, '00]
 - Build a reduced space from multiple eigenvectors
 - Commonly used in recent papers
 - A preferable approach...

Why use multiple eigenvectors?

- **Approximates the optimal cut** [Shi-Malik, '00]
 - Can be used to approximate optimal k -way normalized cut
- **Emphasizes cohesive clusters**
 - Increases the unevenness in the distribution of the data
 - Associations between similar points are amplified, associations between dissimilar points are attenuated
 - The data begins to “approximate a clustering”
- **Well-separated space**
 - Transforms data to a new “embedded space”, consisting of k orthogonal basis vectors
- **Multiple eigenvectors prevent instability due to information loss**

How to select K?

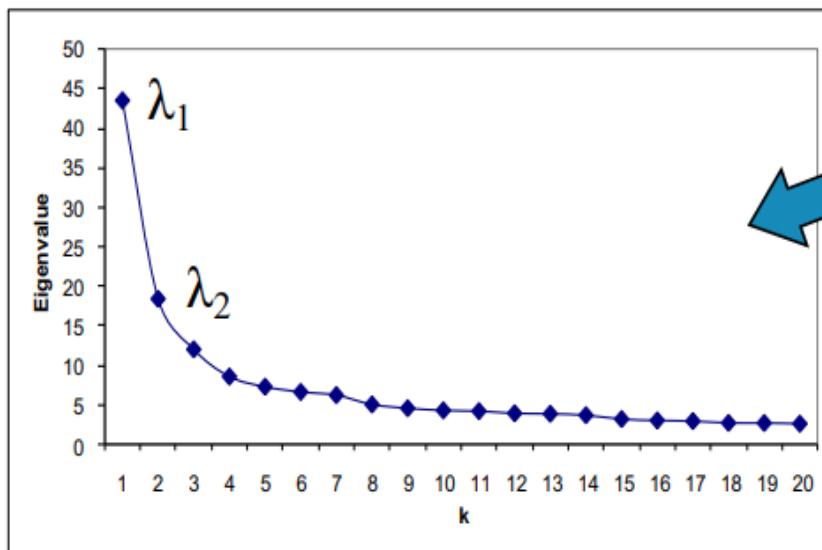
■ Eigengap:

- The difference between two consecutive eigenvalues

■ Most stable clustering is generally given by the value k that maximizes eigengap Δ_k :

$$\Delta_k = |\lambda_k - \lambda_{k-1}|$$

■ Example:



Note eigenvalues
are sorted in
descending order

$$\max \Delta_k = |\lambda_2 - \lambda_1|$$

⇒ Choose
 $k = 2$

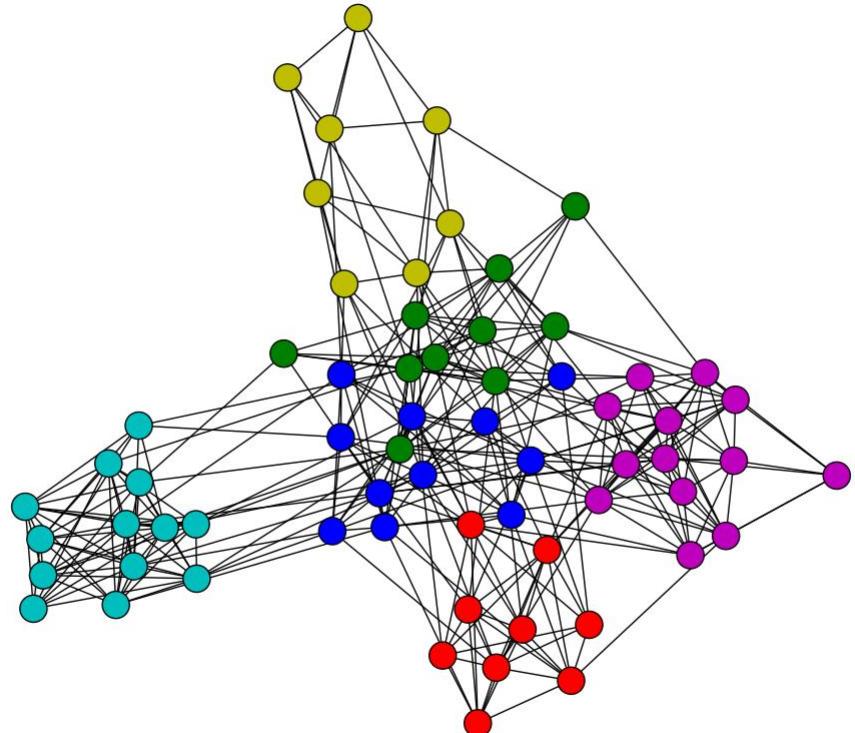
Spectral Clustering: Graph = Matrix

$W^*v_1 = v_2$ “propogates weights from neighbors”

$\mathbf{W} \cdot \mathbf{v} = \lambda \mathbf{v}$: \mathbf{v} is an eigenvector with eigenvalue λ

- smallest eigenvcs of $D-A$ are largest eigenvcs of A
- smallest eigenvcs of $I-W$ are largest eigenvcs of W

Q: How do I pick \mathbf{v} to be an eigenvector for a block-stochastic matrix?

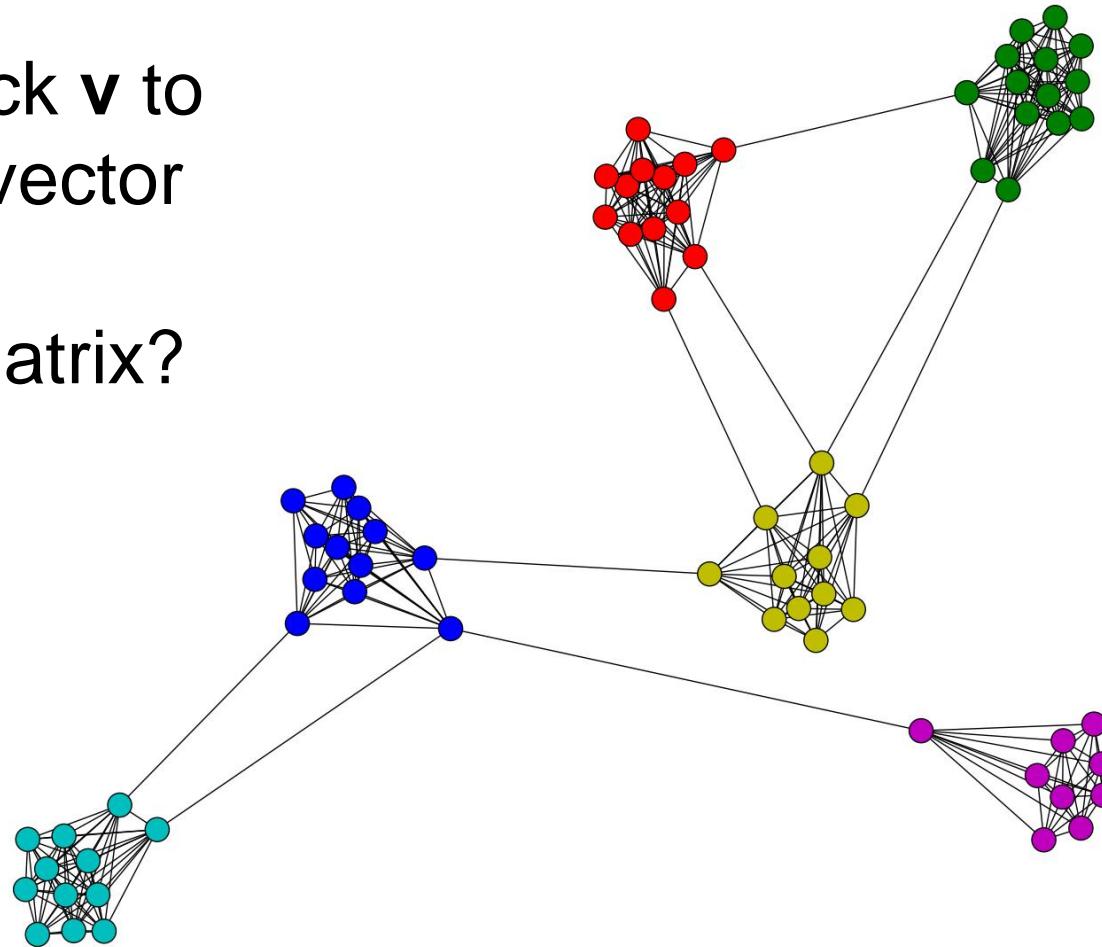


Spectral Clustering: Graph = Matrix

$W^*v_1 = v_2$ “propogates weights from neighbors”

$W \cdot v = \lambda v : v$ is an eigenvector with eigenvalue λ

How do I pick v to
be an eigenvector
for a block-
stochastic matrix?



Spectral Clustering: Graph = Matrix

$W^*v_1 = v_2$ “propogates weights from neighbors”

$\mathbf{W} \cdot \mathbf{v} = \lambda \mathbf{v}$: \mathbf{v} is an eigenvector with eigenvalue λ

- smallest eigenvectors of $D-A$ are largest eigenvectors of A
- smallest eigenvectors of $I-W$ are largest eigenvectors of W

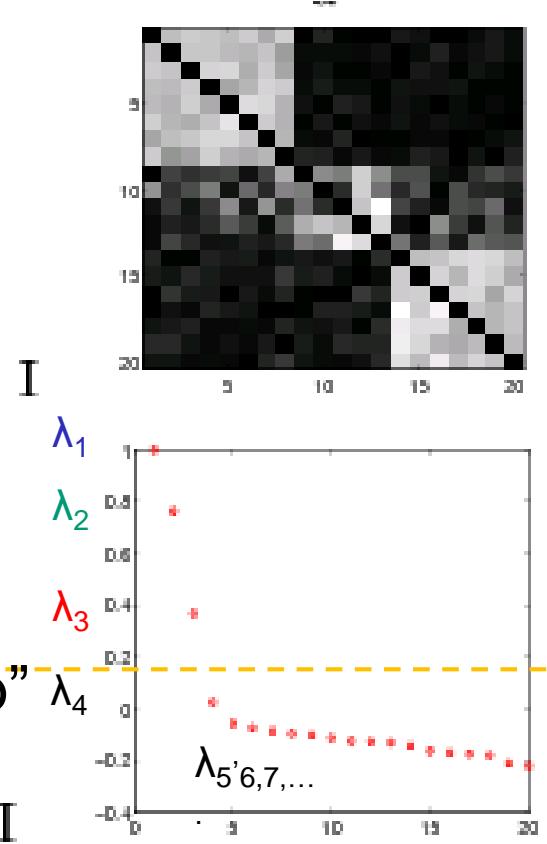
Suppose each $y(i) = +1$ or -1 :

- Then y is a cluster indicator that cuts the nodes into two
- what is $y^T(D-A)y$? The cost of the graph cut defined by y
- what is $y^T(I-W)y$? Also a cost of a graph cut defined by y
- How to minimize it?
 - Turns out: to minimize $y^T X y / (y^T y)$ find *smallest* eigenvector of X
 - But: this will not be $+1/-1$, so it's a “relaxed” solution

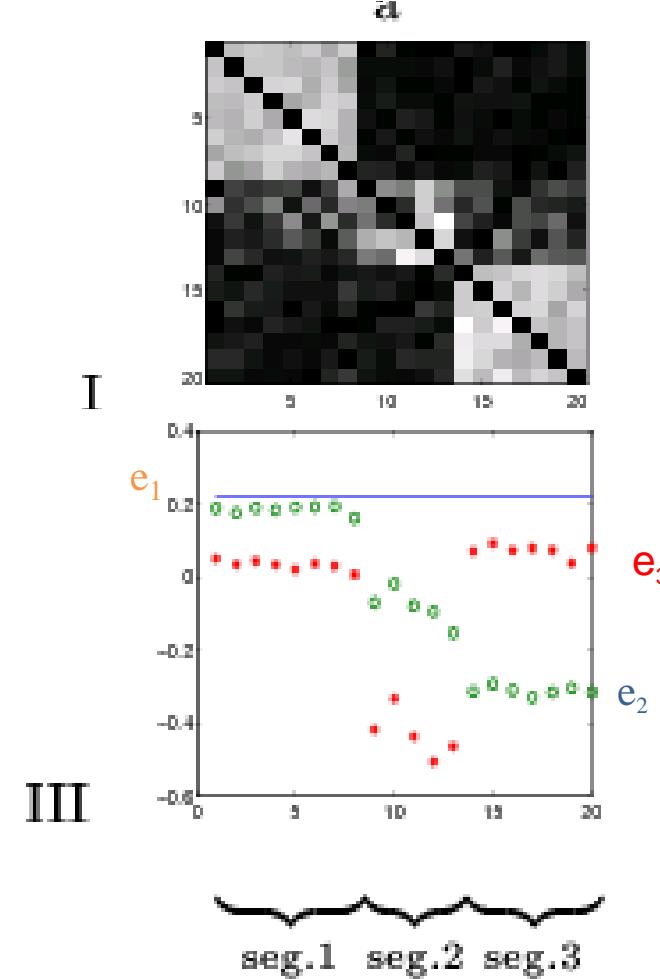
Spectral Clustering: Graph = Matrix

$W^*v_1 = v_2$ “propogates weights from neighbors”

$W \cdot v = \lambda v$: v is an eigenvector with eigenvalue λ



[Shi & Meila, 2002]



seg.1 seg.2 seg.3

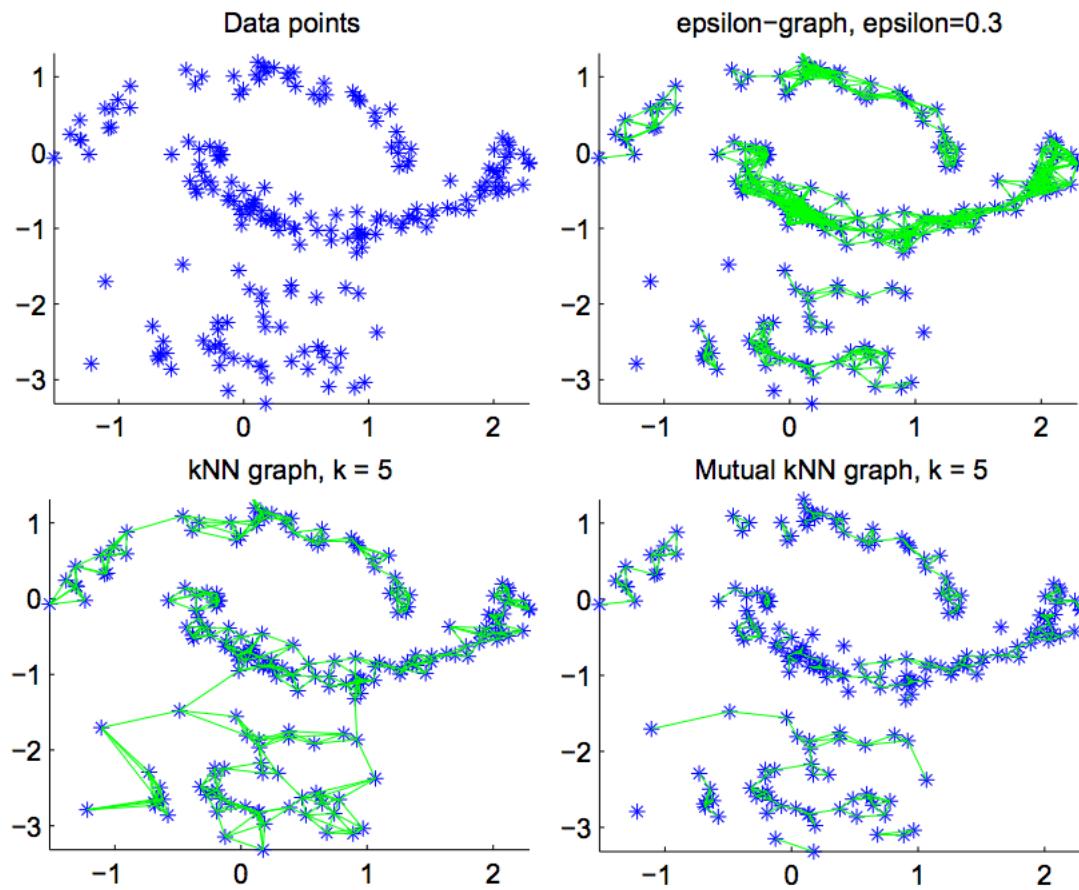


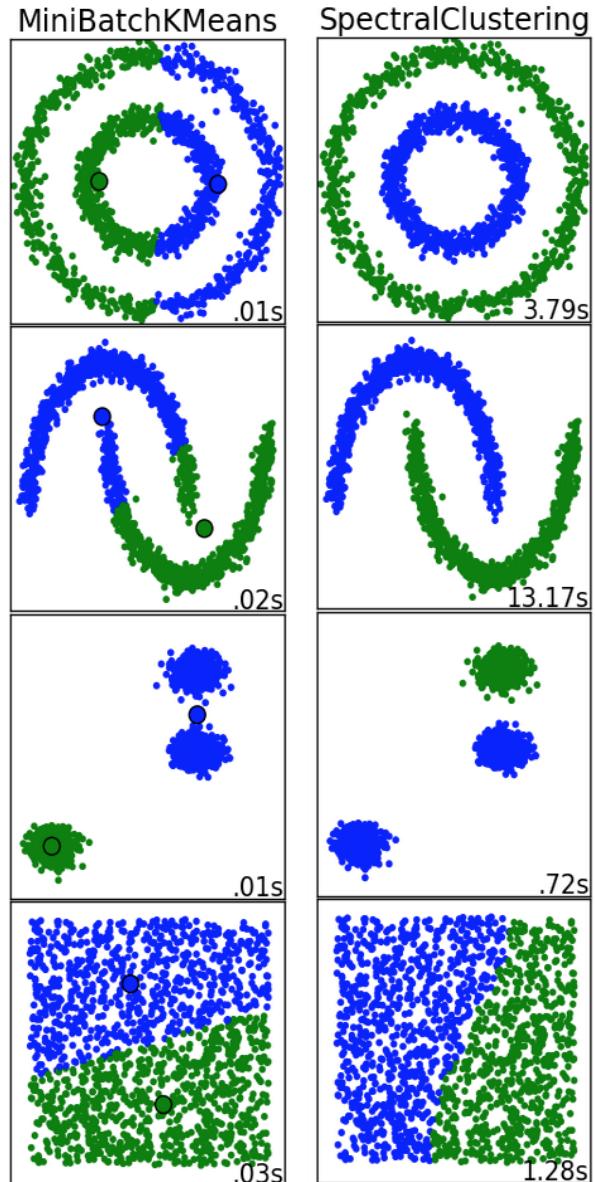
Figure 3: Different similarity graphs, see text for details.

Spectral Clustering: Pros and Cons

- Elegant, and well-founded mathematically
- Works quite well when relations are approximately transitive (like similarity)
- Very noisy datasets cause problems
 - “Informative” eigenvectors need not be in top few
 - Performance can drop suddenly from good to terrible
- Expensive for very large datasets
 - Computing eigenvectors is the bottleneck

Use cases and runtimes

- K-Means
 - FAST
 - “Embarrassingly parallel”
 - Not very useful on anisotropic data
- Spectral clustering
 - Excellent quality under many different data forms
 - Much slower than K-Means



Further Reading

- Spectral Clustering Tutorial:

http://www.informatik.uni-hamburg.de/ML/contents/people/luxburg/publications/Luxburg07_tutorial.pdf